

Operating Systems

CS2006

Lecture 6

Inter-Process Communication (IPC)

15th February 2023

Dr. Rana Asif Rehman

Introduction

- A process has access to the memory which constitutes its own address space.
- When a child process is created, the only way to communicate between a parent and a child process is:
 - The parent receives the exit status of the child
- So far, we've discussed communication mechanisms only during process creation/termination
- Processes may need to communicate during their life time.

Cooperating Processes

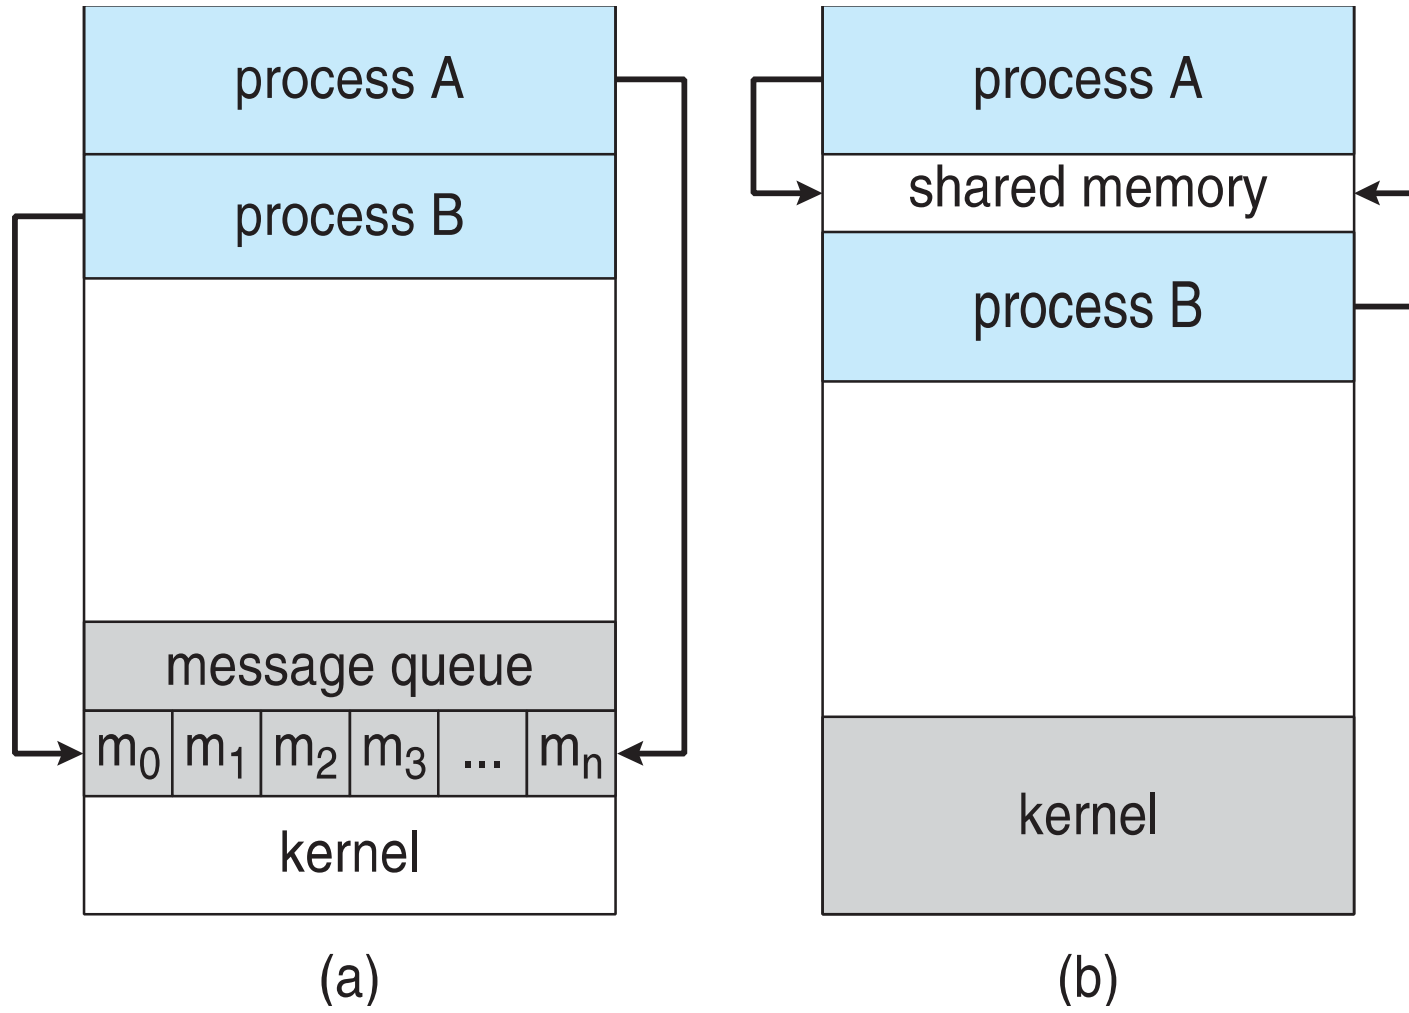
- **Independent** process cannot affect or be affected by the execution of another process.
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience
- Dangers of process cooperation
 - Data corruption, deadlocks, increased complexity
 - Requires processes to synchronize their processing

Inter-Process Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions.
- Two main types
 - Shared Memory
 - Message Passing

Communications Models

(a) Message passing. (b) shared memory.



Interprocess Communication – Shared Memory

- ❑ An area of memory shared among the processes that wish to communicate
- ❑ The communication is under the control of the users processes not the operating system.
- ❑ Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- ❑ Synchronization is discussed in great details in Chapter 6.

Producer-Consumer Problem

- ❑ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
- ❑ To allow producer and consumer processes to run concurrently, we must have available a buffer of items that can be filled by the producer and emptied by the consumer.
- ❑ This buffer will reside in a region of memory that is shared by the producer and consumer processes.
- ❑ A producer can produce one item while the consumer is consuming another item.
- ❑ The producer and consumer must be synchronized, so that the consumer does not try to consume an item that has not yet been produced.

Producer-Consumer Problem

Two types of buffers can be used.

- **unbounded-buffer** places no practical limit on the size of the buffer
- **bounded-buffer** assumes that there is a fixed buffer size
- Let's look more closely at how the bounded buffer illustrates inter- process communication using shared memory....

Bounded-Buffer – Shared-Memory Solution

- The following variables reside in a region of memory shared by the producer and consumer processes:
- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Lets code this....

Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* buffer is full - do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Bounded Buffer – Consumer

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* buffer empty - do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next consumed */  
}
```

- This scheme allows at most $\text{BUFFER_SIZE} - 1$ items in the buffer at the same time.

Message Passing (IPC)

- **Message system** — processes communicate with each other **without** resorting to **shared variables**.
- IPC facility provides two operations:
 - **send**(message) — message size fixed or variable
 - **receive**(message)
- If P and Q wish to communicate, they need to:
 - establish a communication link between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (i.e., shared memory, hardware bus)
 - logical (direct/indirect, blocking/non-blocking, automatic/explicit buffering)

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Message Passing Systems

- Exchange messages over a communication link
- Methods for implementing the communication link and primitives (send/receive):
 1. Direct or Indirect communications (Naming)
 2. Symmetric or Asymmetric communications (blocking versus non-blocking)
 3. Buffering

Direct Communication

- Processes must **name each other** explicitly:
 - **send** (P, message) – send a message to process P
 - **receive**(Q, message) – receive a message from process Q
- Properties of communication link
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - Processes need to know each other's identity
 - The link may be unidirectional, but is usually bi-directional.

Disadvantage: a process must know the name or ID of the process(es) it wishes to communicate with

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports).
 - Each mailbox has a unique id.
 - Processes can communicate only if they share a mailbox.
- Properties of communication link
 - Link established only if processes have a shared mailbox
 - A link may be associated with many processes.
 - Each pair of processes may share several communication links.
 - Link may be unidirectional or bi-directional.

Indirect Communication (Cont.)

- Operations provided by the OS
 - **create** a new **mailbox**
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - **send**(A, message) – send a message to mailbox A
 - **receive**(A, message) – receive a message from mailbox A

Indirect Communication (Cont.)

- Mailbox sharing
 - P1, P2, and P3 share mailbox A.
 - P1, sends; P2 and P3 receive.
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing can be blocking or non-blocking
- Four types of synchronization
 - Blocking send:
 - sender blocked until message received by mailbox or process
 - Nonblocking send:
 - sender resumes operation immediately after sending
 - Blocking receive:
 - receiver blocks until a message is available
 - Nonblocking receive:
 - receiver returns immediately with either a valid or null message.

Buffering

- Messages exchanged by processes reside in temporary queue
- Three ways to implement queues

1. **Zero capacity**

No messages may be queued within the link, requires sender to block until receiver retrieves message.

2. **Bounded capacity**

Link has finite number of message buffers. If no buffers are available then sender must block until one is freed up.

3. **Unbounded capacity**

Link has unlimited buffer space, consequently send never needs to block.

Communication in Client/Server Systems

- Sockets
- Remote procedure calls
- Pipes
- What types of Inter-Process Communication might each of these use?

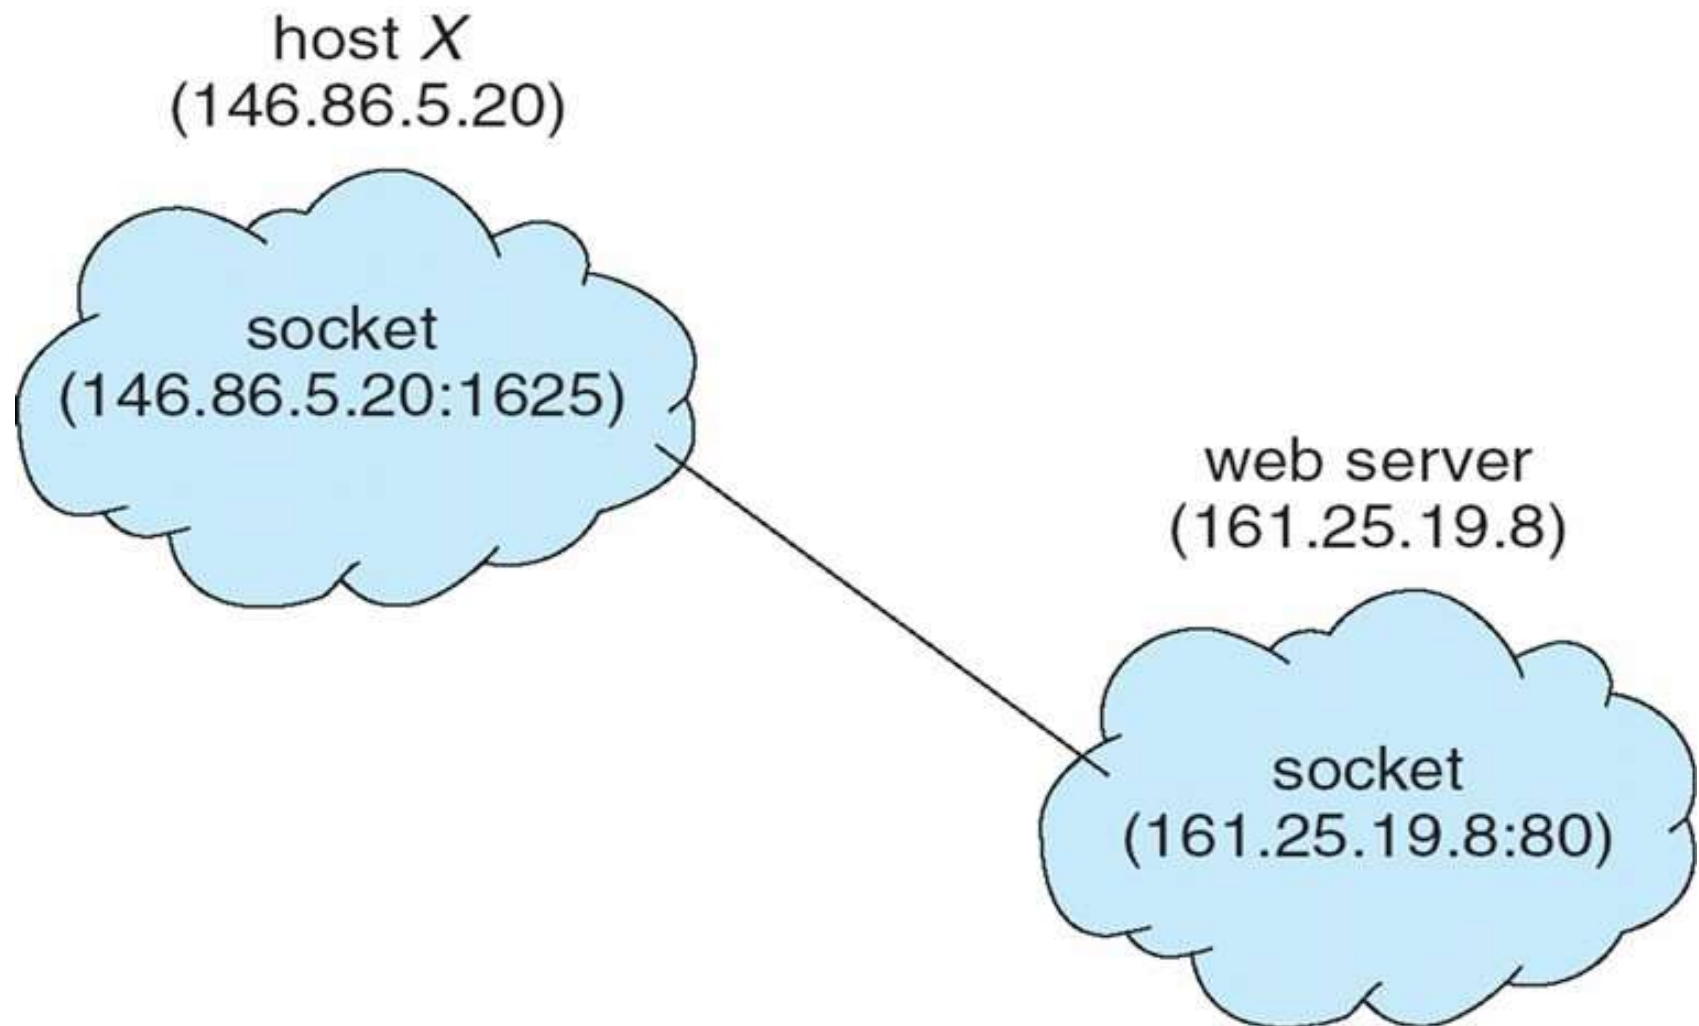
1. Sockets

- Defined as a point of communication
- A pair of process communicating over a network employ a pair of sockets
- Socket = IP Address : Port Number
- All ports below 1024 are considered well known
 - FTP: 21
 - Telnet: 23
 - Web server: 80

Sockets

- Client initiates a requests for connection, it is assigned a port by host computer
- Connection on one host must be unique
- Sockets can be
 - Connection-oriented (TCP) socket
 - Connectionless (UDP) socket
- Considered Low level mode of communication
 - Shares unstructured stream of bytes

Sockets



Sockets in Java

□ Three types of sockets

- **Connection-oriented (TCP)**

- **Connectionless (UDP)**

- **MulticastSocket**
class— data can be sent to
multiple recipients

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

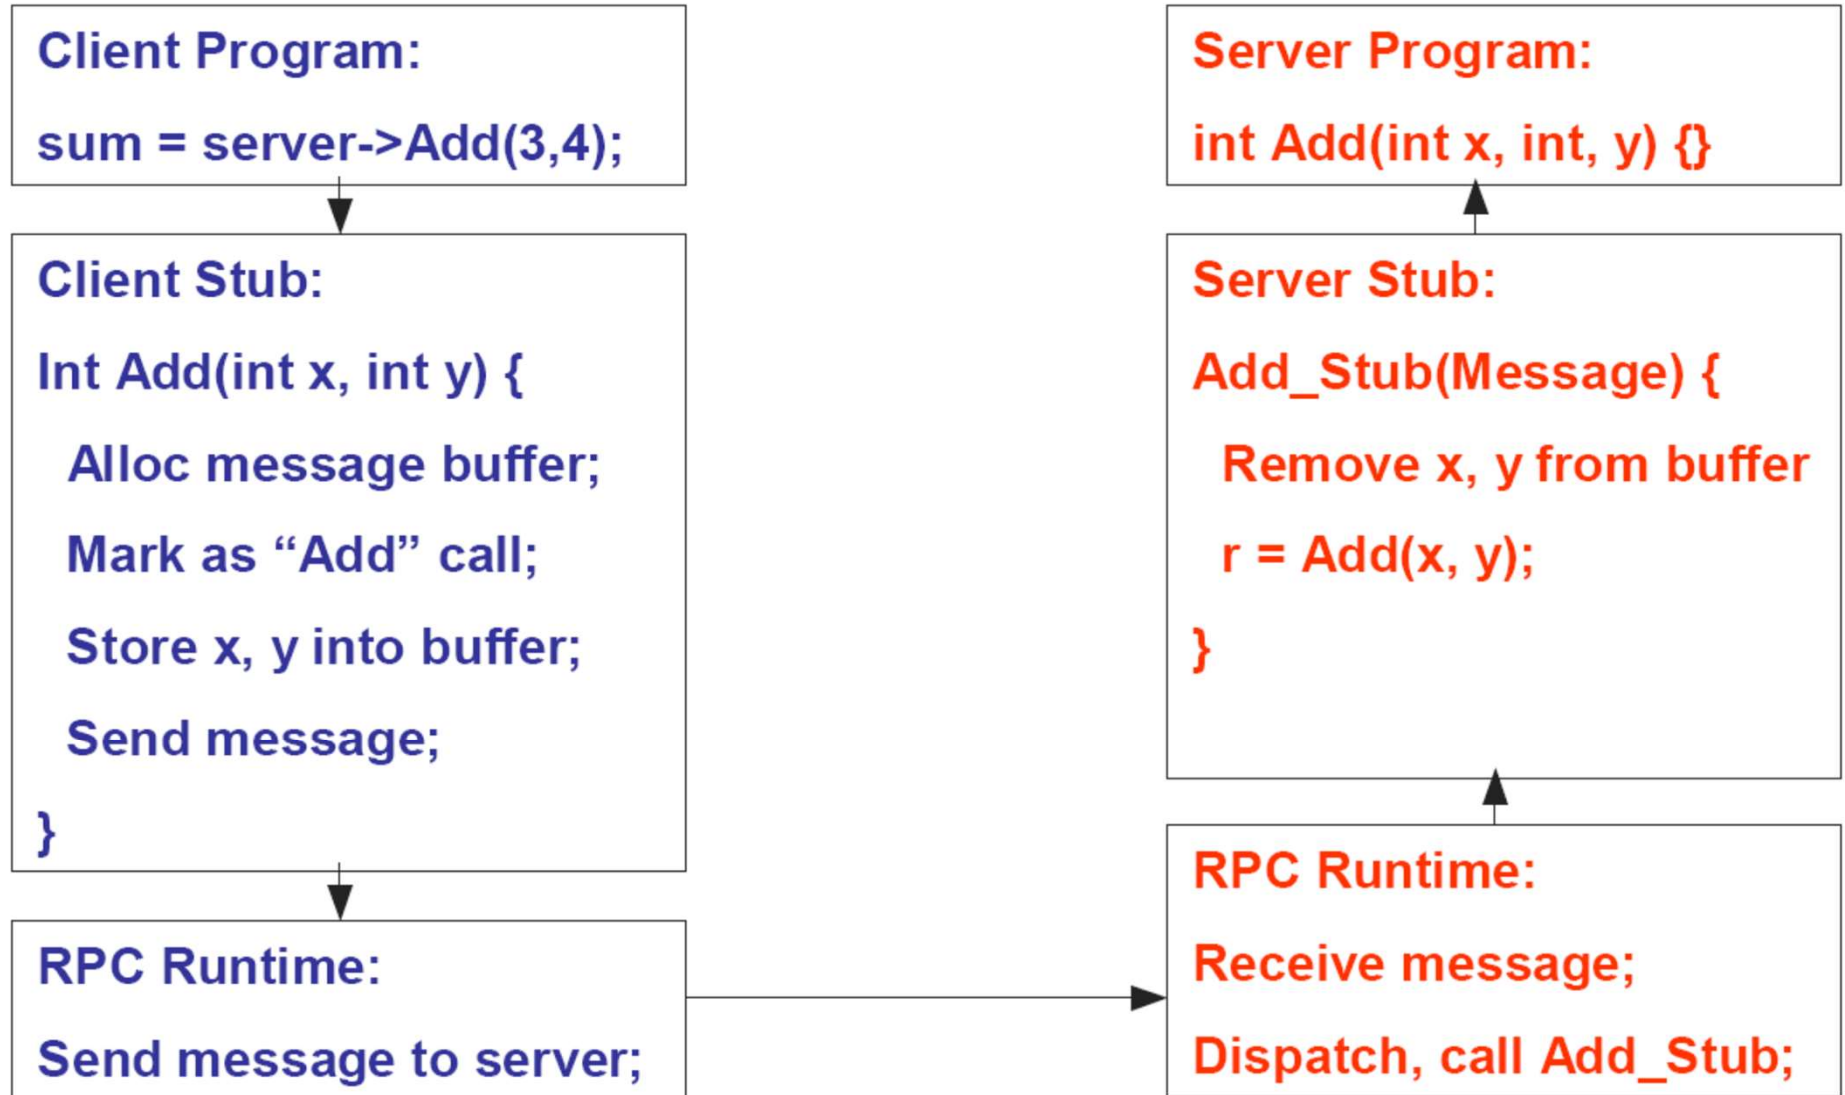
                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

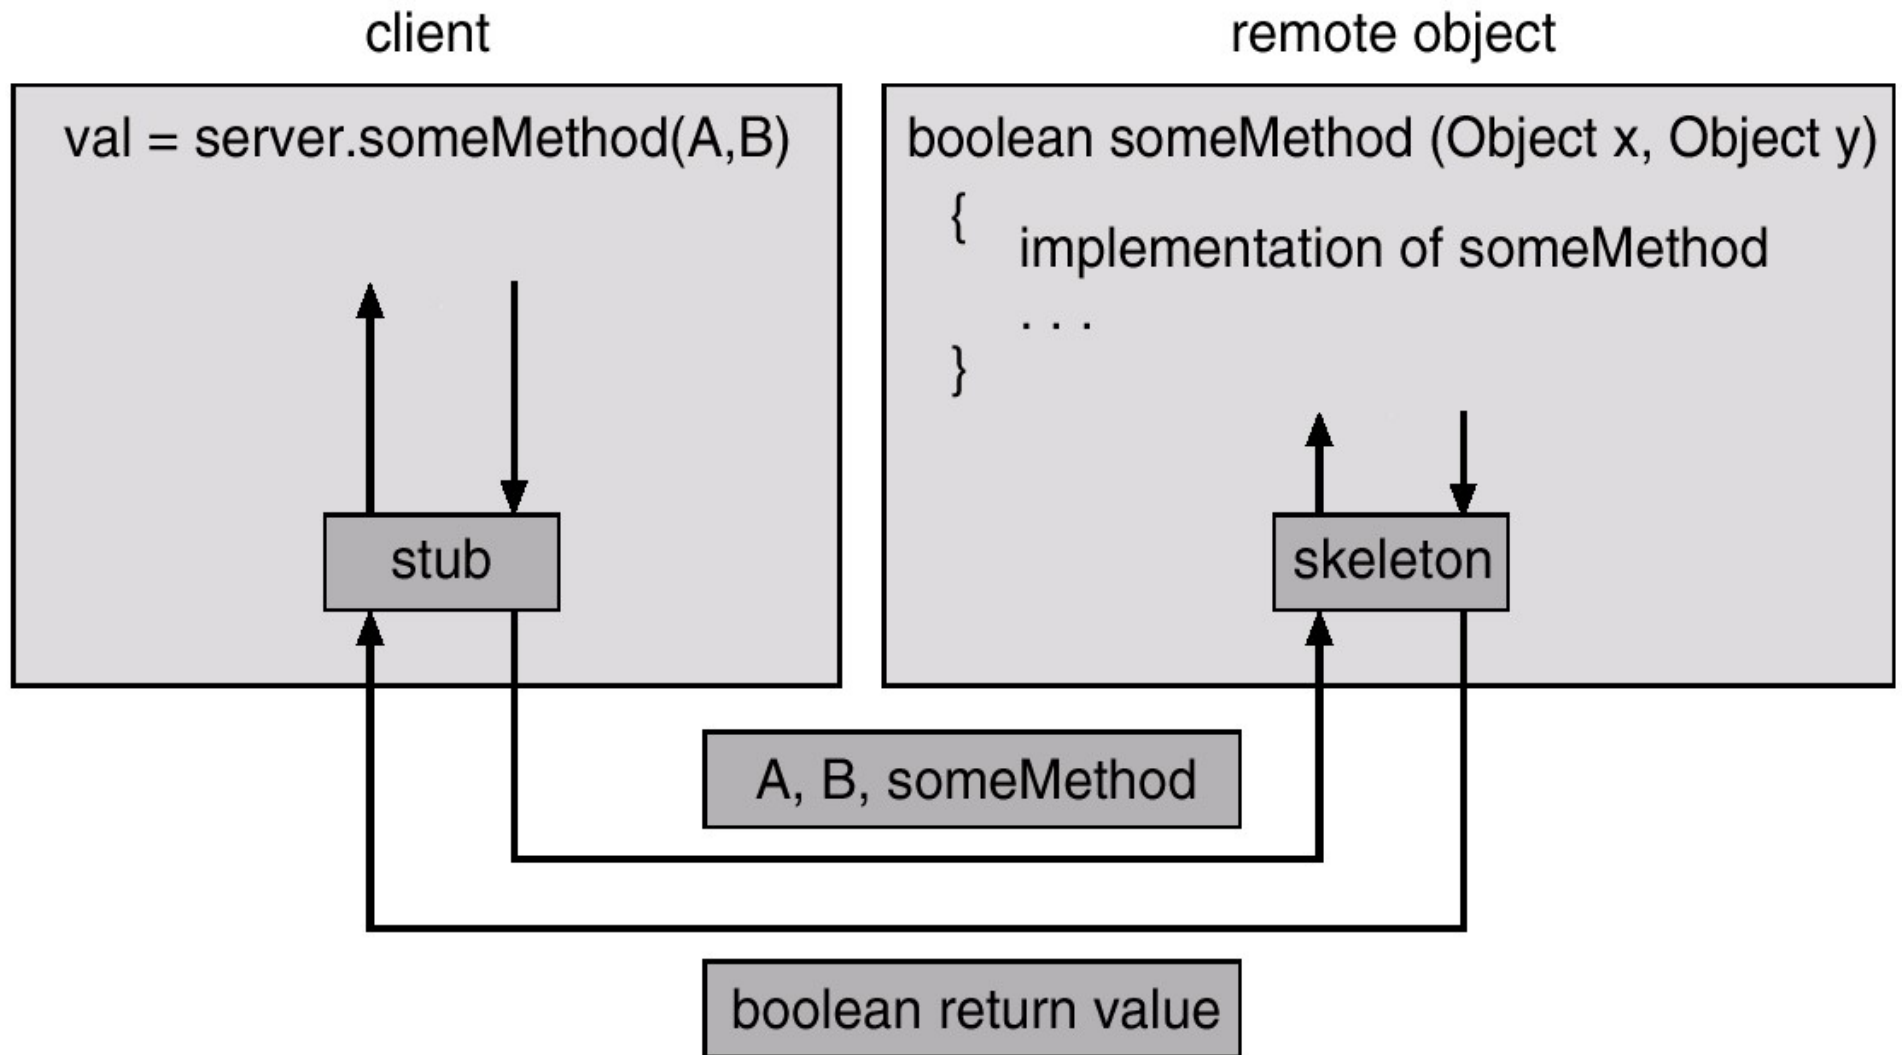
2. Remote Procedure Calls (RPC)

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and marshals the parameters.
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.

RPC: An Example



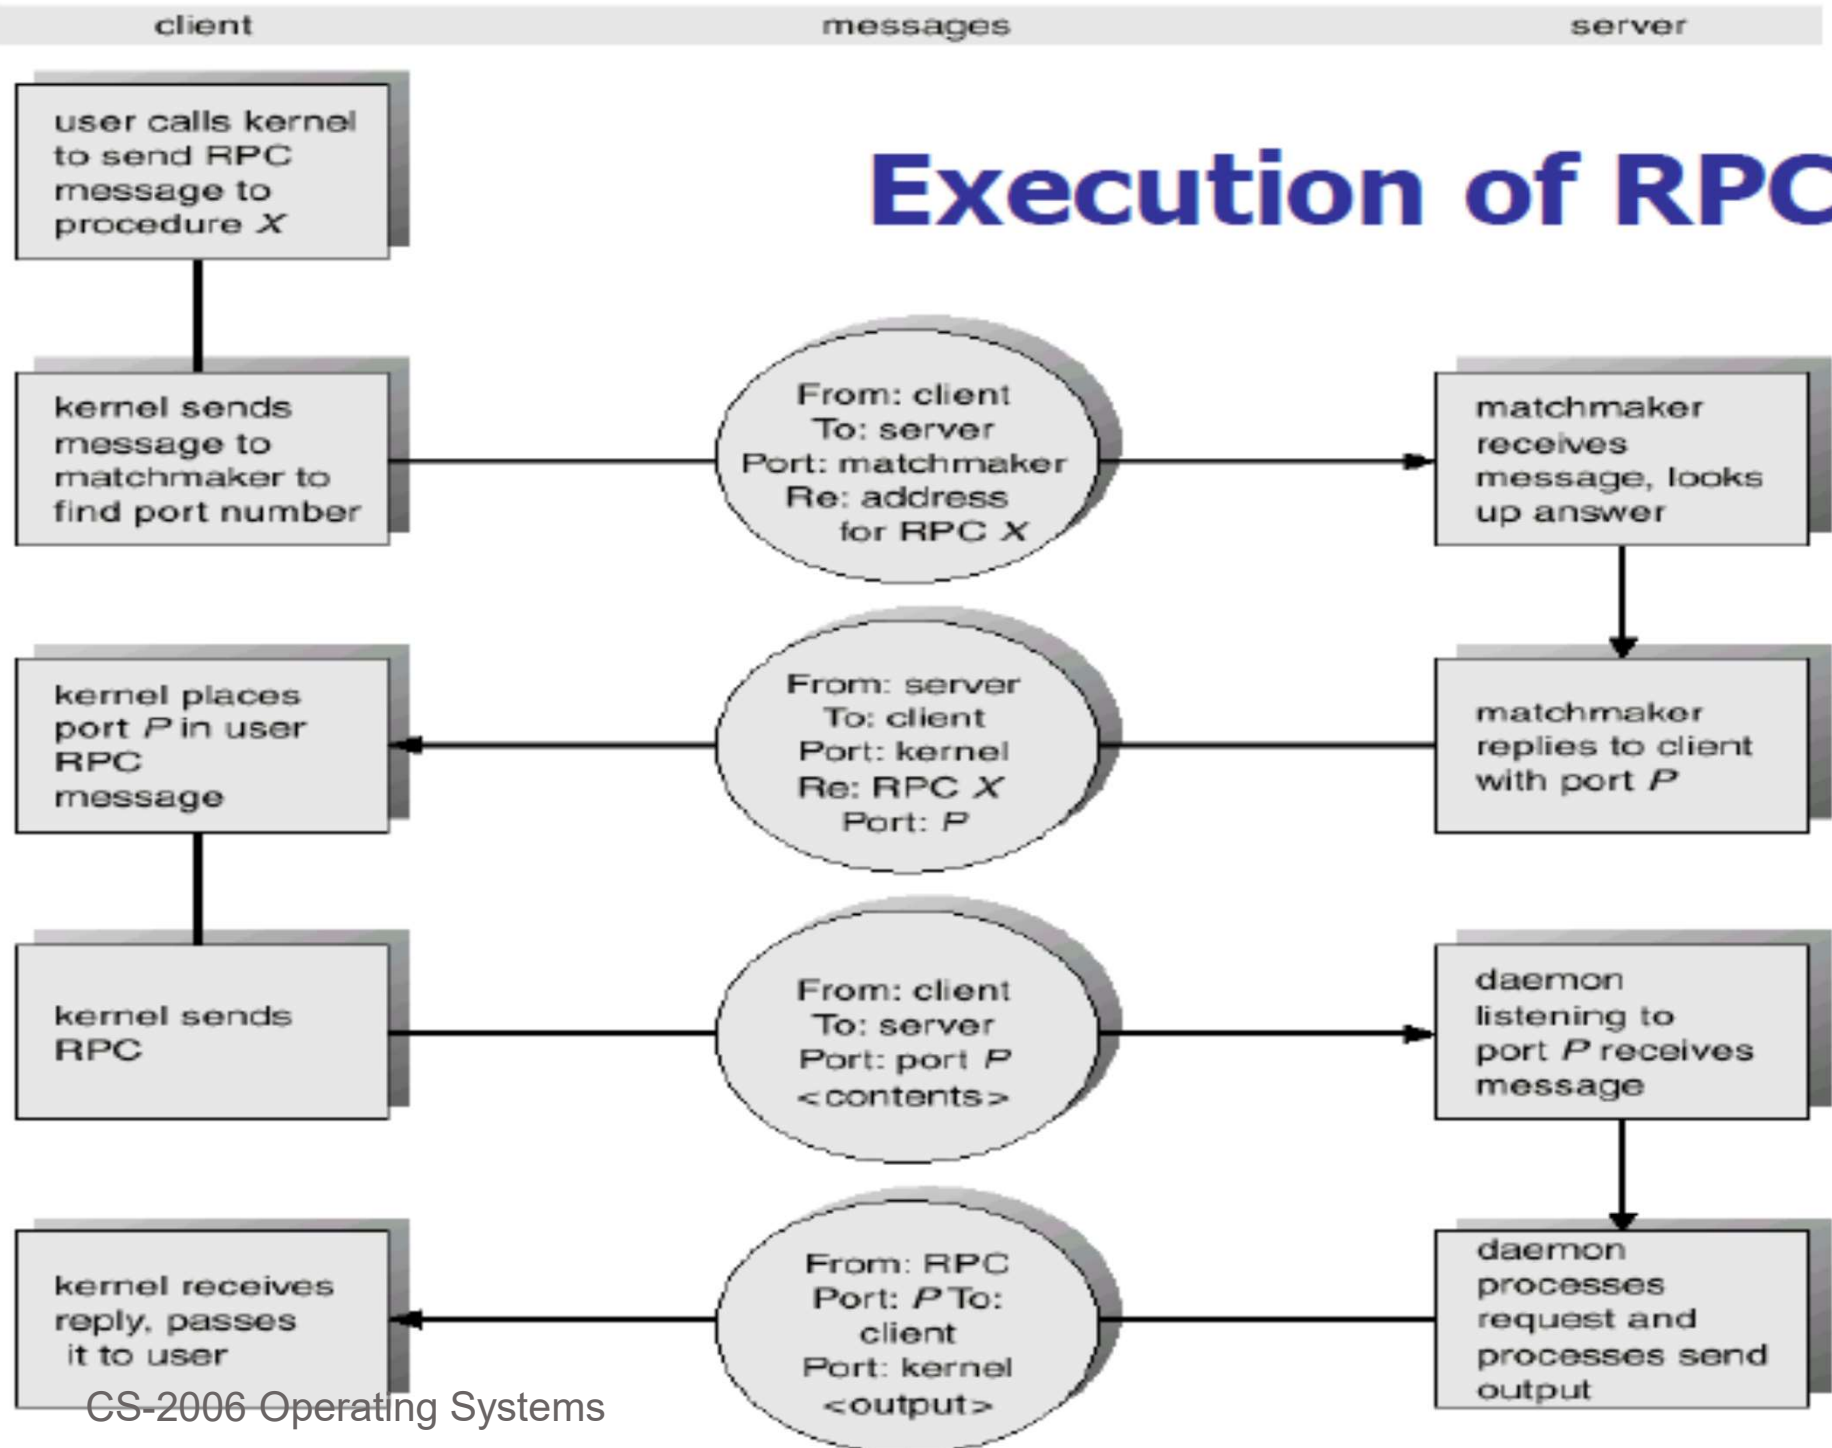
Marshalling Parameters



Remote Procedure Calls

- Problem exists in data representation
 - Big endian
 - Little endian
- Many RPC systems define a machine independent representation
 - External Data Representation (XDR)
- Semantics of a call
 - Messages are acted on *exactly once* rather than *at-most once*
- Port binding of RPC server
 - Fixed at compile time
 - Done dynamically by **rendezvous** mechanism (*match maker*)


Execution of RPC



3. Pipes

- IPC mechanism in early UNIX systems
- Four issues must be considered
 - Unidirectional or bi-directional communication
 - If bidirectional then is it half duplex or full duplex
 - Must a relationship (such as *parent-child*) exist between the communicating processes?
 - Can pipe communicate over a network?

Pipes: Shared info in kernel's memory



Pipe Creation (Unix)

```
#include <unistd.h>
```

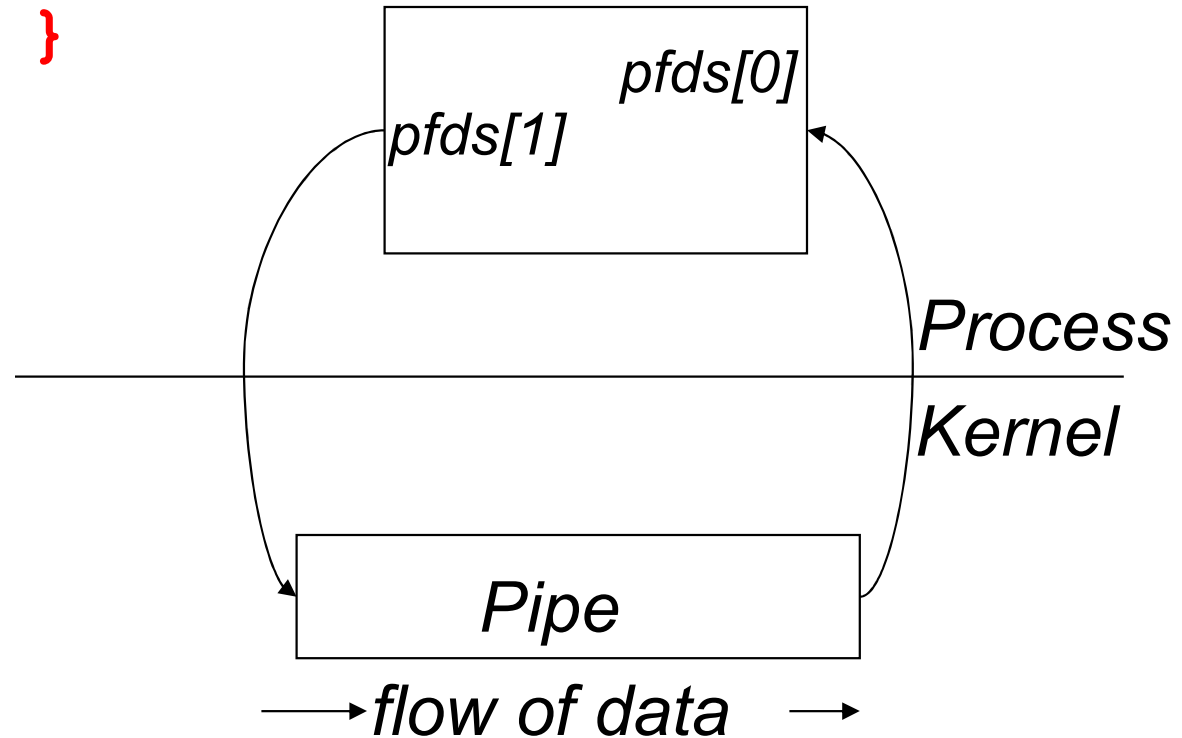
```
int pipe(int filedes[2]);
```

- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by filedes
 - `filedes[0]` is for reading
 - `filedes[1]` is for writing.
- On success, zero is returned.
- On error, -1 is returned

Pipe Creation

```
int main()
{  int pfd[2];

  if (pipe(pfd) == -1)
    {  perror("pipe");
        exit(1);  }
}
```



Reading/Writing from/to a Pipe

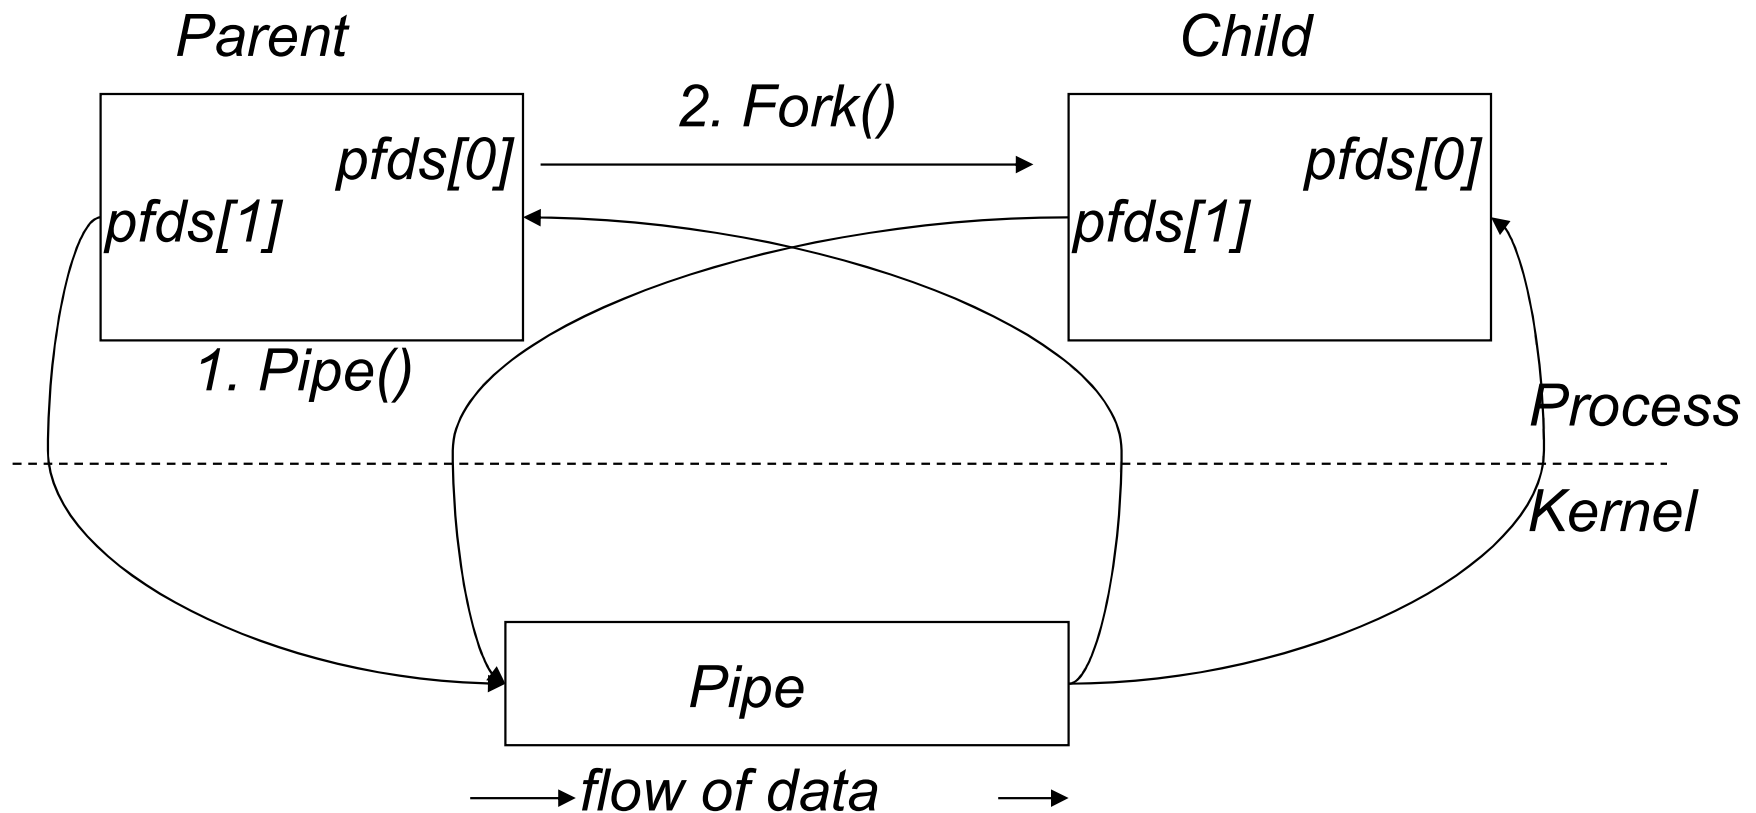
- `int read(int filedescriptor, char *buffer, int bytetoread) ;`
- `int write(int filedescriptor, char *buffer, int bytetowrite) ;`

Example

```
int main()
{ int pfds[2];
  char buf[30];
  if (pipe(pfds) == -1) {
      perror("pipe");
      exit(1); }
  printf("writing to file descriptor #%d\n", pfds[1]);
  write(1, "test", 5);?

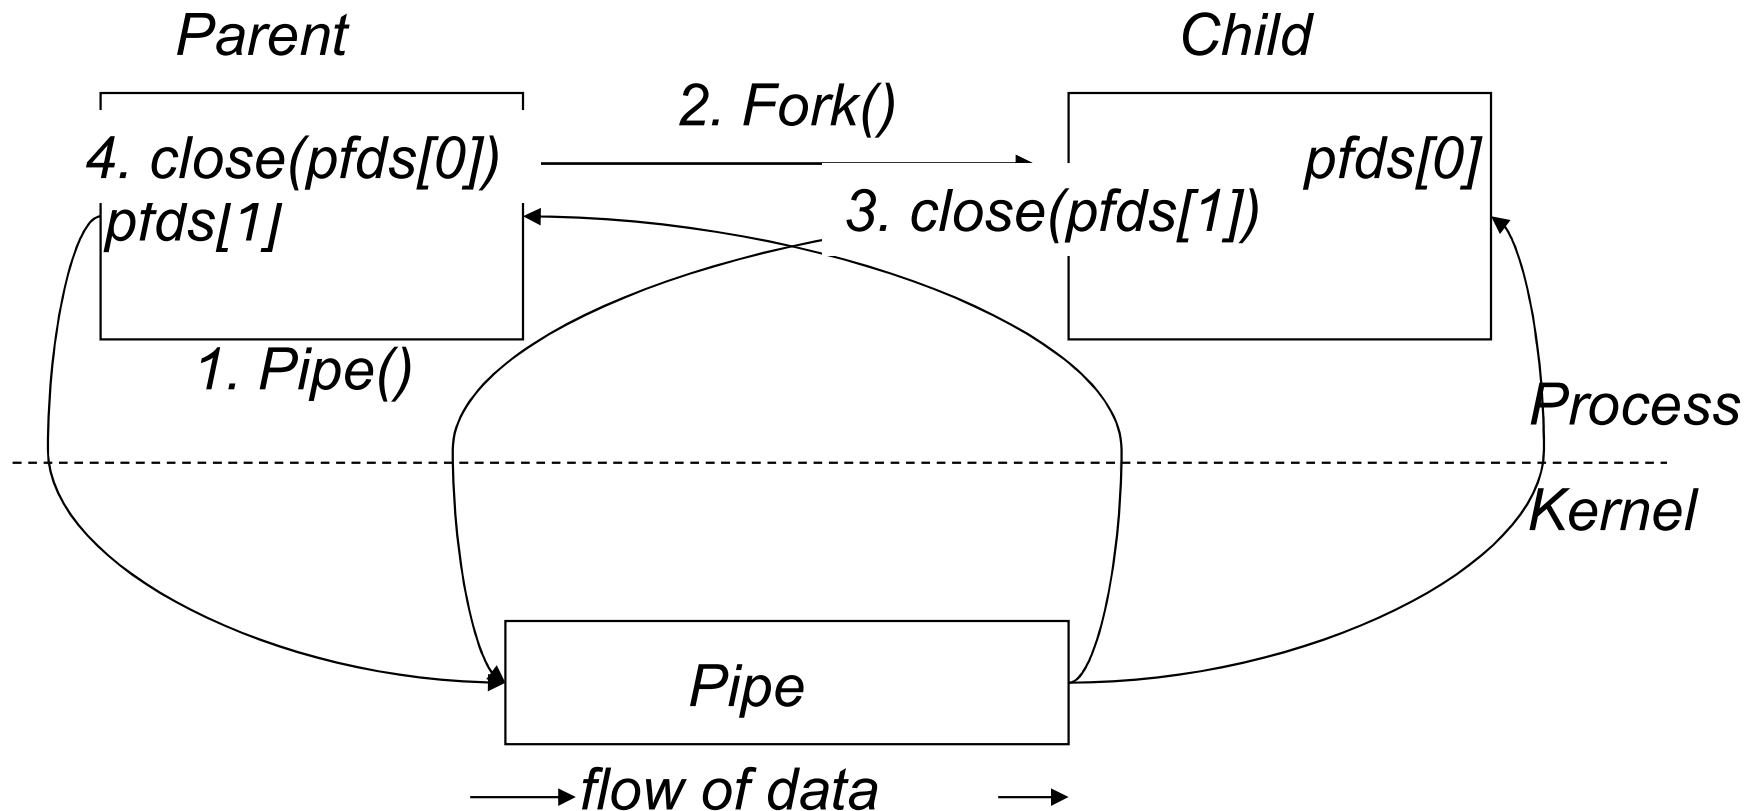
  printf("reading from file descriptor
  #0\n");
  read(0, buf, 5);?
  printf("read %s\n", buf);
}
```

A Channel between two processes



A Channel between two processes

- To allow one way communication each process should close one end of the pipe.




Pipes

- Two types of pipes
 - Ordinary pipes
 - Named pipes

Ordinary pipes

- **Ordinary pipes** allow communication in standard producer-consumer form
- **Producer** writes to one end of the pipe (the *write-end*)
- **Consumer** reads from other end of the pipe (the *read-end* of the pipe)
- Ordinary pipes are therefore unidirectional
- Required parent-child relationship between communicating processes
- In windows known as *Anonymous pipes*

Ordinary pipe



Named pipes

- Named pipes are more powerful than ordinary pipes
- Communication is bidirectional
- UNIX provides half duplex while Windows provides full duplex communication
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows
- Called FIFO in UNIX
- In UNIX, communicating process reside on the same machine (socket is required for intermachine comm.)
- In Windows, communicating process may reside on either the same machine or different machines

References

- Operating System Concepts (Silberschatz, 9th edition)
Chapter 3