

Operating Systems

CS2006

Lecture 7

CPU Scheduling

20th February 2023

Dr. Rana Asif Rehman

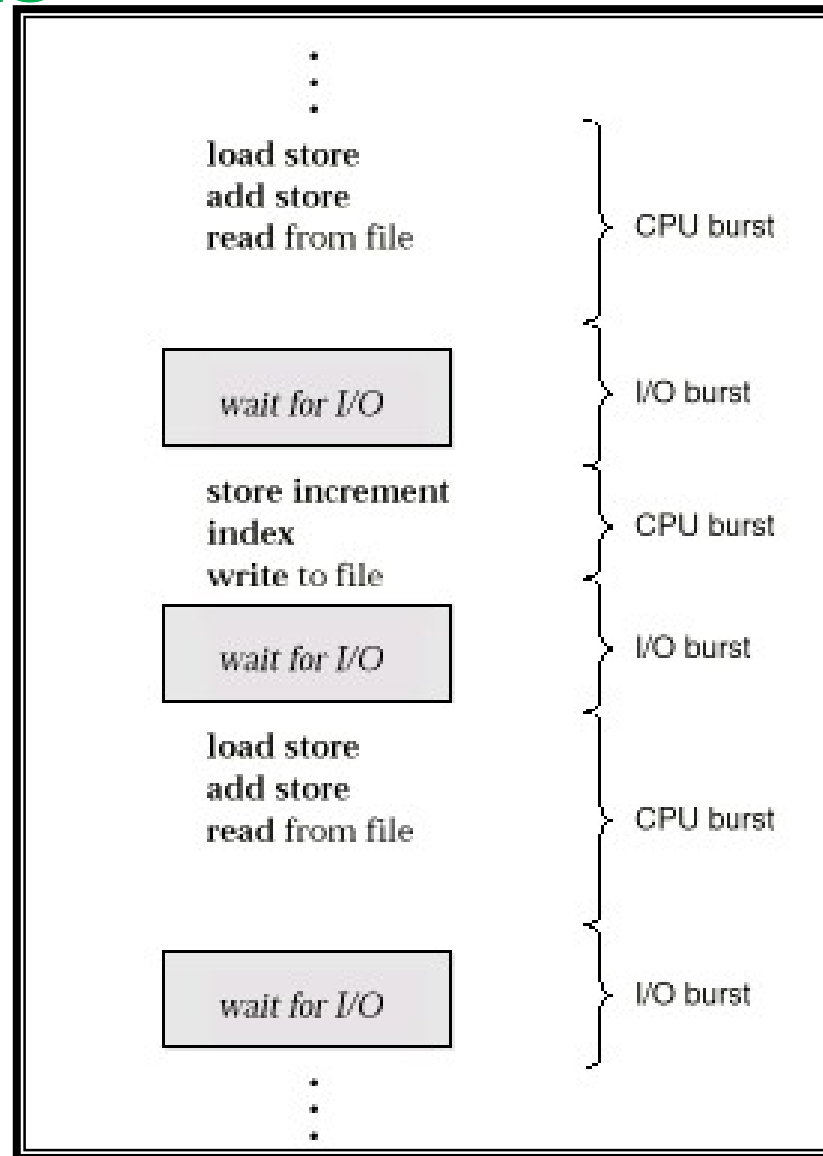
What's in today's lecture

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms

Basic Concepts

- In multiprogramming systems, where there is more than one process runnable (ready), the OS must decide which one to run next
- We already talked about the **dispatcher**, whose job is to allow the next process to run
- It's intimately associated with a **scheduler** whose job is to decide (using a **scheduling algorithm**) which process to have the dispatcher dispatch
- We will focus on the scheduling, understand the problems associated with it, and look at some scheduling algorithms

Alternating Sequence of CPU And I/O Bursts



Nature of Process

- Not all processes have an even mix of CPU and I/O usage
- A number crunching program may do a lot of computation and minimal I/O
 - This is an example of a **CPU-BOUND** process
- A data processing job may do little computation and a lot of I/O
 - This is an example of an **I/O-BOUND** process

Schedulers

- Chooses a ready process and activates it
- Policy – which process to choose
- Choice of policy depends on goals:
 - Maximize throughput – avoid idle time and overhead
 - Fairness – all jobs get same service
 - Minimum average waiting time – good service to some
 - Good service to “important” jobs

Schedulers

- Short term Scheduler or CPU Scheduling
- Long term Scheduler or Job Scheduler
- Medium Term Scheduler

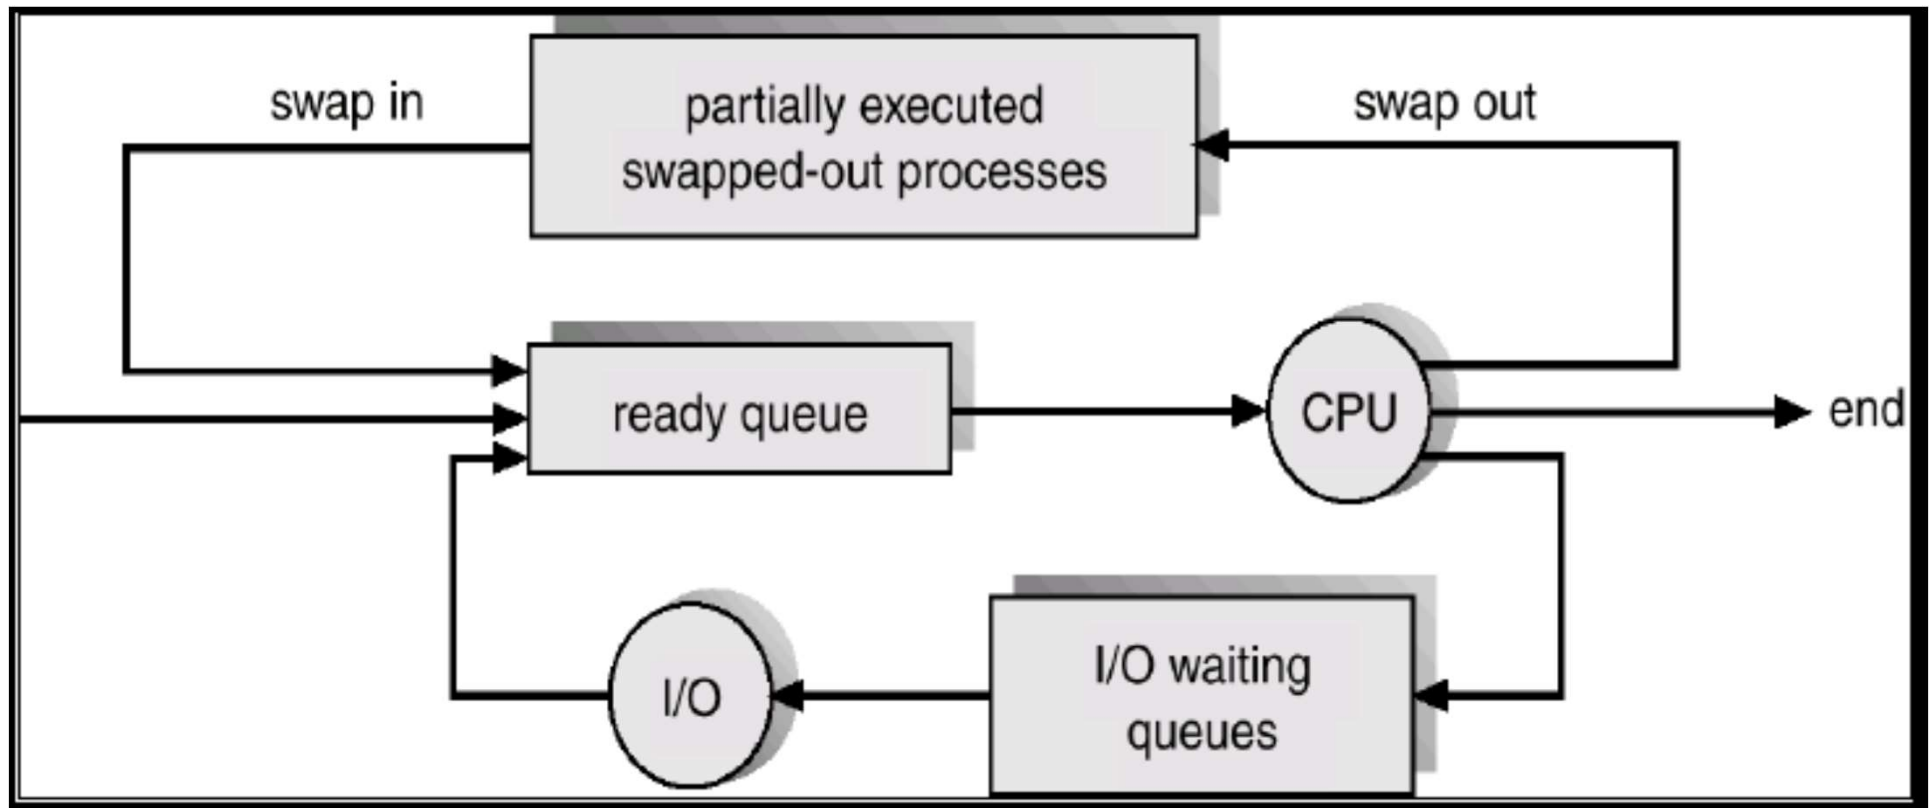
Schedulers

- Short term Scheduler or CPU Scheduling
 - Which program is to be run next
- Long term Scheduler or Job Scheduler
 - Which ready jobs should be brought to memory
 - May need to invoke only when a process leaves the system
 - Must make a careful selection

Schedulers

- Sometimes OS may swap a blocked process to disk to free up more memory
- Or to improve process mix
- This is also called **Swapping**
- This scheduler responsible for this task is known as **Intermediate or Medium Term Scheduler**

Addition of Medium Term Scheduling



Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow Design Objective? **(must be fast)**
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow **(may be slower)**
- The long-term scheduler controls the degree of *multiprogramming*

Scheduling

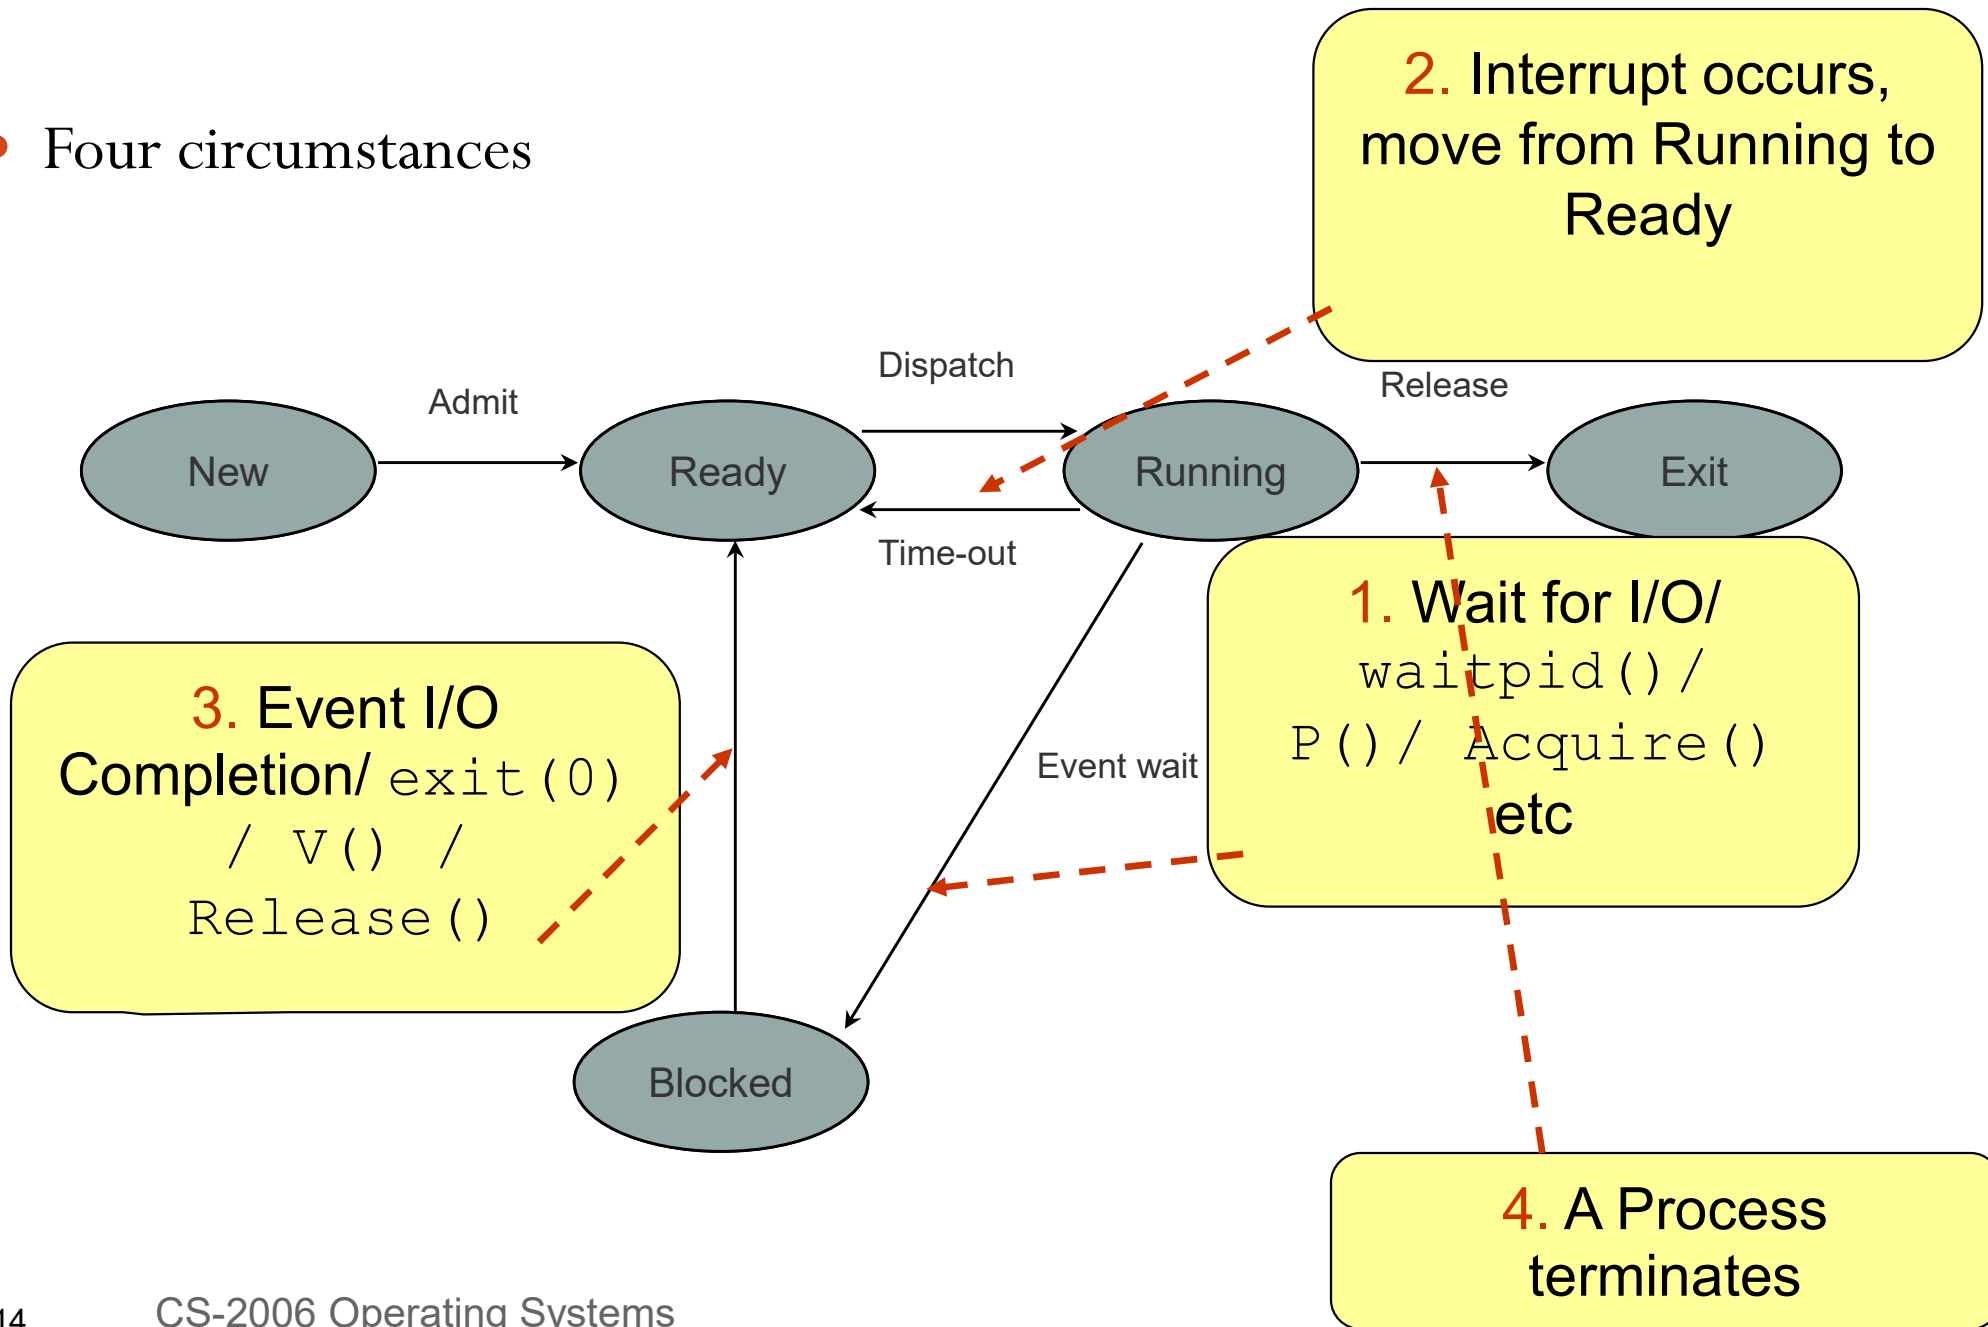
- Short term scheduler (CPU Scheduler)
 - Whenever the CPU becomes idle, a process must be selected for execution
 - The Process is selected from the Ready queue
- Ready queue is not necessarily a FIFO queue
- It can be
 - Priority based
 - A Tree
 - Unordered linked list etc

Short term Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is *nonpreemptive*
- All other scheduling is *preemptive*

When to select a new process to Run

- Four circumstances

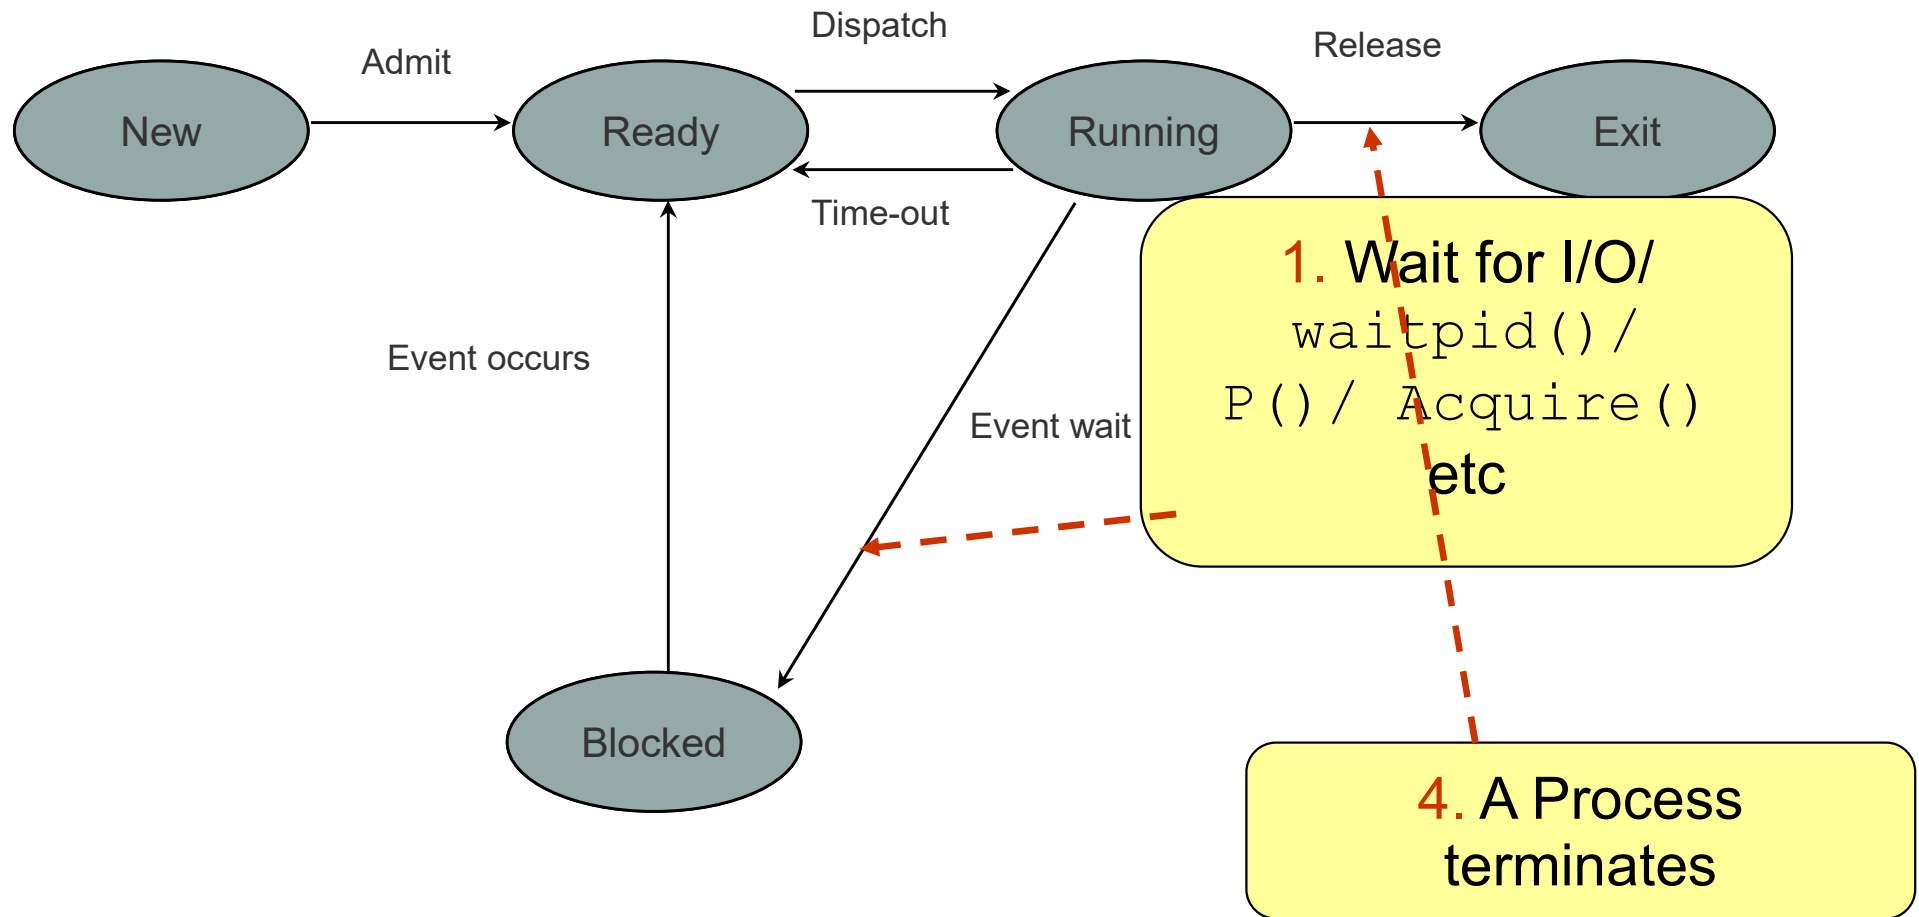


Non Preemptive Scheduling

- Once the CPU has been allocated to a process
- The process keeps it until
 - It Terminates
 - Or has to wait for:
 - I/O
 - Mutex
 - Child process
 - Semaphore
 - Conditional Variables etc
- There is no way, to get the CPU back, FORCEFULLY

Non Preemptive Scheduling

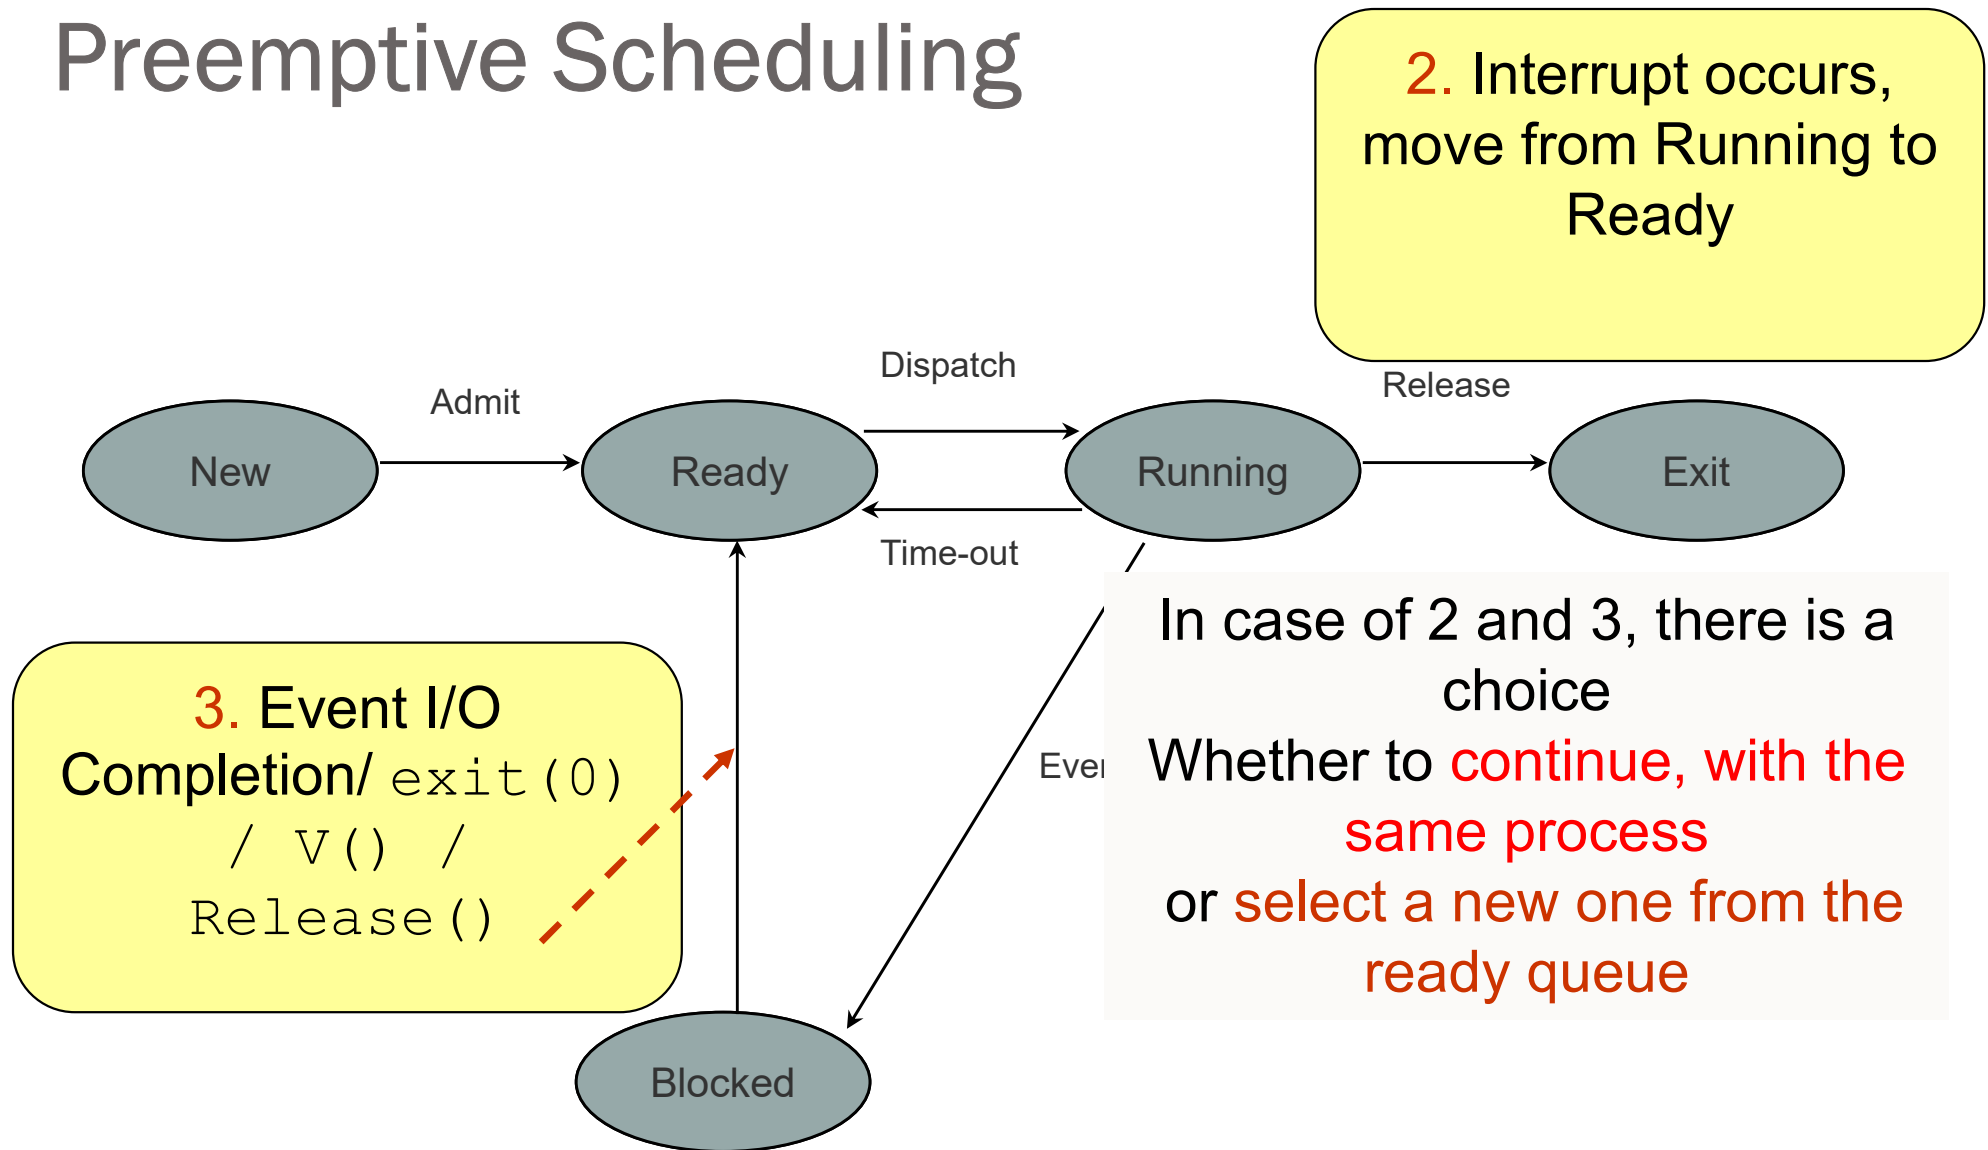
- Only the case 1 and 4
- Must select a new process, if any, from the Ready Queue



Preemptive Scheduling

- Possible to get the CPU back from a process, even if the process has not completed its execution.

Preemptive Scheduling



Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - Switching context
 - Switching to user mode
 - Jumping to the proper location in the user program to restart that program
- *Dispatch latency* — time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

- CPU Utilization
 - Keep the CPU as busy as is possible
 - May range from 0% to 100%
- Throughput
 - Number of processes completed per unit time
 - E.g. long processes
 - 1 process / hr
 - Short processes
 - 10 processes / hr

Scheduling Criteria

- Turnaround Time

- How long it take to execute a Process

$$\text{Turnaround} = \text{Completion_Time} \\ - \text{Submission_Time}$$

$$\text{Turnaround} = \text{Wait_Time}_{\text{GetIntoMemory}} \\ + \text{Wait_Time}_{\text{ReadyQueue}} \\ + \text{Wait_Time}_{\text{BlockQueue}} \\ + \text{CPU_Execution_Time}$$

Scheduling Criteria

- Scheduling Algorithm does not effect the waiting time in Block Queue
- It only effect the Waiting Time in the Ready Queue
- Waiting Time
 - Sum of the periods spent waiting in the Ready Queue

$$W.T = \text{Finish Time} - \text{Burst Time} - \text{Arrival Time}$$

Scheduling Criteria

- Turnaround Time is not a good criteria for Interactive Systems
- A process may
 - Produce “Some” output
 - Computes new results, while previous results are output to the user
- Response Time
- **Response_Time =
First_Response_Start_Time
– Submission_Time**

Additional Scheduling Criteria

- There are also other elements to consider:
 - **Priority/Importance** of work — hopefully more important work can be done first
 - **Fairness** — hopefully eventually everybody is served
 - Implement policies to increase priority as we wait longer... (this is known as “priority aging”)
 - **Deadlines** — some processes may have hard or soft deadlines that must be met
 - Consistency and/or predictability may be a factor as well, especially in interactive systems

Optimization Criteria

- We would like to **Maximize**
 - CPU Utilization
 - Throughput
- And **Minimize**
 - Turnaround Time
 - Waiting Time
 - Response Time

Scheduling Algorithms

- First come, First serve (FCFS)
- Shortest Job First (SJF)
- Priority Scheduling
- Round-Robin Scheduling
- Multi-level Queue Scheduling
- Multi-level Feed back queue Scheduling


1. First-Come, First-Served (FCFS) Scheduling

- Simplest scheduling algorithm:
 - Run jobs in order that they arrive
- Uni-programming:
 - Run until done
- Multi-programming:
 - Run until done or Blocks on I/O
- Nonpreemptive
 - A Process keeps CPU until done or I/O
- Advantage:
 - Simplicity

FCFS Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)


- Suppose that the processes arrive in the order P_2 , P_3 , P_1
- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case

Convoy Effect

process	service time t_s	turnaround time t_t	waiting time t_w
A	10	10	0
B	1	11	10
C	3	14	11
D	4	18	14
AVERAGE		13.25	8.75



A long CPU-bound job may hog the CPU and force shorter (or I/O-bound) jobs to wait for prolonged periods. This in turn may lead to a lengthy queue of ready jobs, and hence to the **'convoy effect'**

FCFS Scheduling (Cont.)

- FCFS is:
 - Non-preemptive
 - Ready queue is a FIFO queue
 - Jobs arriving are placed at the end of ready queue
 - First job in ready queue runs to completion of CPU burst
- **Advantages:** simple, low overhead
- **Disadvantages:** long waiting time, inappropriate for interactive systems, large fluctuations in average turnaround time are possible

References

- Operating System Concepts (Silberschatz, 9th edition)
Chapter 5