

Operating Systems

CS2006

Lecture 8

CPU Scheduling-II

22nd February 2023

Dr. Rana Asif Rehman

What's in today's lecture

- Basic Concepts
- Scheduling Criteria
- **Scheduling Algorithms**

Scheduling Algorithms

- First come, First serve (FCFS)
- Shortest Job First (SJF)
- Priority Scheduling
- Round-Robin Scheduling
- Multi-level Queue Scheduling
- Multi-level Feed back queue Scheduling


2. Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - **Nonpreemptive** — once CPU is given to the process it cannot be preempted until the process completes its CPU burst
 - **Preemptive** — if a new process arrives with CPU burst length less than the remaining time of current executing process, preempt. This scheme is known as the **Shortest-Remaining-Time-First (SRTF)**
- SJF is optimal — gives minimum average waiting time for a given set of processes

Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$


SRTF - Shortest Remaining Time First

- Preemptive version of SJF
- Ready queue ordered on length of time till completion (shortest first)
- Arriving jobs inserted at proper position
- Shortest job
 - Runs to completion (i.e. CPU burst finishes) or
 - Runs until a job with a shorter remaining time arrives (i.e. placed in the ready queue)

Example of Preemptive SJF (i.e., SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- Preemptive SJF (i.e., SRTF)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Shortest-Job-First (SJF) Scheduling

- Ready queue treated as a priority queue based on smallest CPU-time requirement
 - Arriving jobs inserted at proper position in queue
 - Shortest job (1st in queue) runs to completion
- In general, SJF is often used in long-term scheduling
- Advantages: provably optimal w.r.t. average waiting time
- Disadvantages: Unimplementable at the level of short-term CPU scheduling. Also, starvation is possible!
- Can do it approximately: use exponential averaging to predict length of next CPU burst
==> pick shortest predicted burst next!

Determining Length of Next CPU Burst

- Can only estimate the length
- Can be done by using the length of previous CPU bursts, using exponential averaging

1. t_n = actual length of n^{th} CPU burst

2. τ_{n+1} = predicted value for the next (i.e., $n^{th} + 1$) CPU burst

3. τ_n = predicted value for the n^{th} CPU burst

4. $\alpha, 0 \leq \alpha \leq 1$

5. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.


$\alpha = 0$ implies making no use of recent history ($\tau_{n+1} = \tau_n$)

$\alpha = 1$ implies $\tau_{n+1} = t_n$ (past prediction not used)

$\alpha = 1/2$ implies weighted (older bursts get less and less weight)

Prediction of the Length of the Next CPU Burst

This figure is for $\alpha = 0.5$ and $\tau_0 = 10$



3. Priority Scheduling


- A priority number (integer) is associated with each process
 - Priority can be internally computed (e.g., may involve time limits, memory usage) or externally (e.g., user type, funds being paid)
 - In SJF, priority is simply the predicted next CPU burst time
- The CPU is allocated to the process with the highest priority (smallest integer might mean highest priority)
- **Starvation** is a problem, where low priority processes may never execute
- **Solution:** as time progresses, the priority of the long waiting (starved) processes is increased. This is called **aging**

Priority scheduling

- Priority scheduling can be **Preemptive** or **Non-Preemptive**
- When a process arrives and enters the Ready Queue
- Its priority is compared with the currently Running Process
- If Higher
 - Preemptive Scheduling
 - Run the New process
 - Non-Preemptive Scheduling
 - Continue running the process

Priority Scheduling

process	priority	service time t_s	turnaround time t_t	waiting time t_w
A	4	10	18	8
B	3	1	8	7
C	2	3	7	4
D	1	4	4	0
AVERAGE			9.25	4.75




4. Round Robin (RR)

- RR reduces the penalty that short jobs suffer with FCFS by preempting running jobs periodically
- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds
- The CPU blocks the current job when its reserved *time quantum (time-slice)* is exhausted
 - The current job is then put at the end of the ready queue if it has not yet completed
 - If the current job is completed, it will exit the system (terminate)

Round Robin (RR)

- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n - 1)q$ time units
- Performance: the critical issue with the RR policy is the length of the quantum q
 - q is large: RR will behave like FIFO and hence interactive processes will suffer
 - q is small: the CPU will be spending more time on context switching
 - q must be large with respect to context switch, otherwise overhead is too high


Time Quantum and Context Switch Time



Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:



- Typically, higher average turnaround than SJF but better *response*

Round Robin's Disadvantage


- Good for Varying sized jobs
- But what about same-sized jobs?
- Assume 2 jobs of time = 100 each:



- Avg completion time?
- $(200 + 200) / 2 = 200$
- How does this compare with FCFS for same two jobs?
- $(100 + 200) / 2 = 150$

Turnaround Time Varies With The Time Quantum

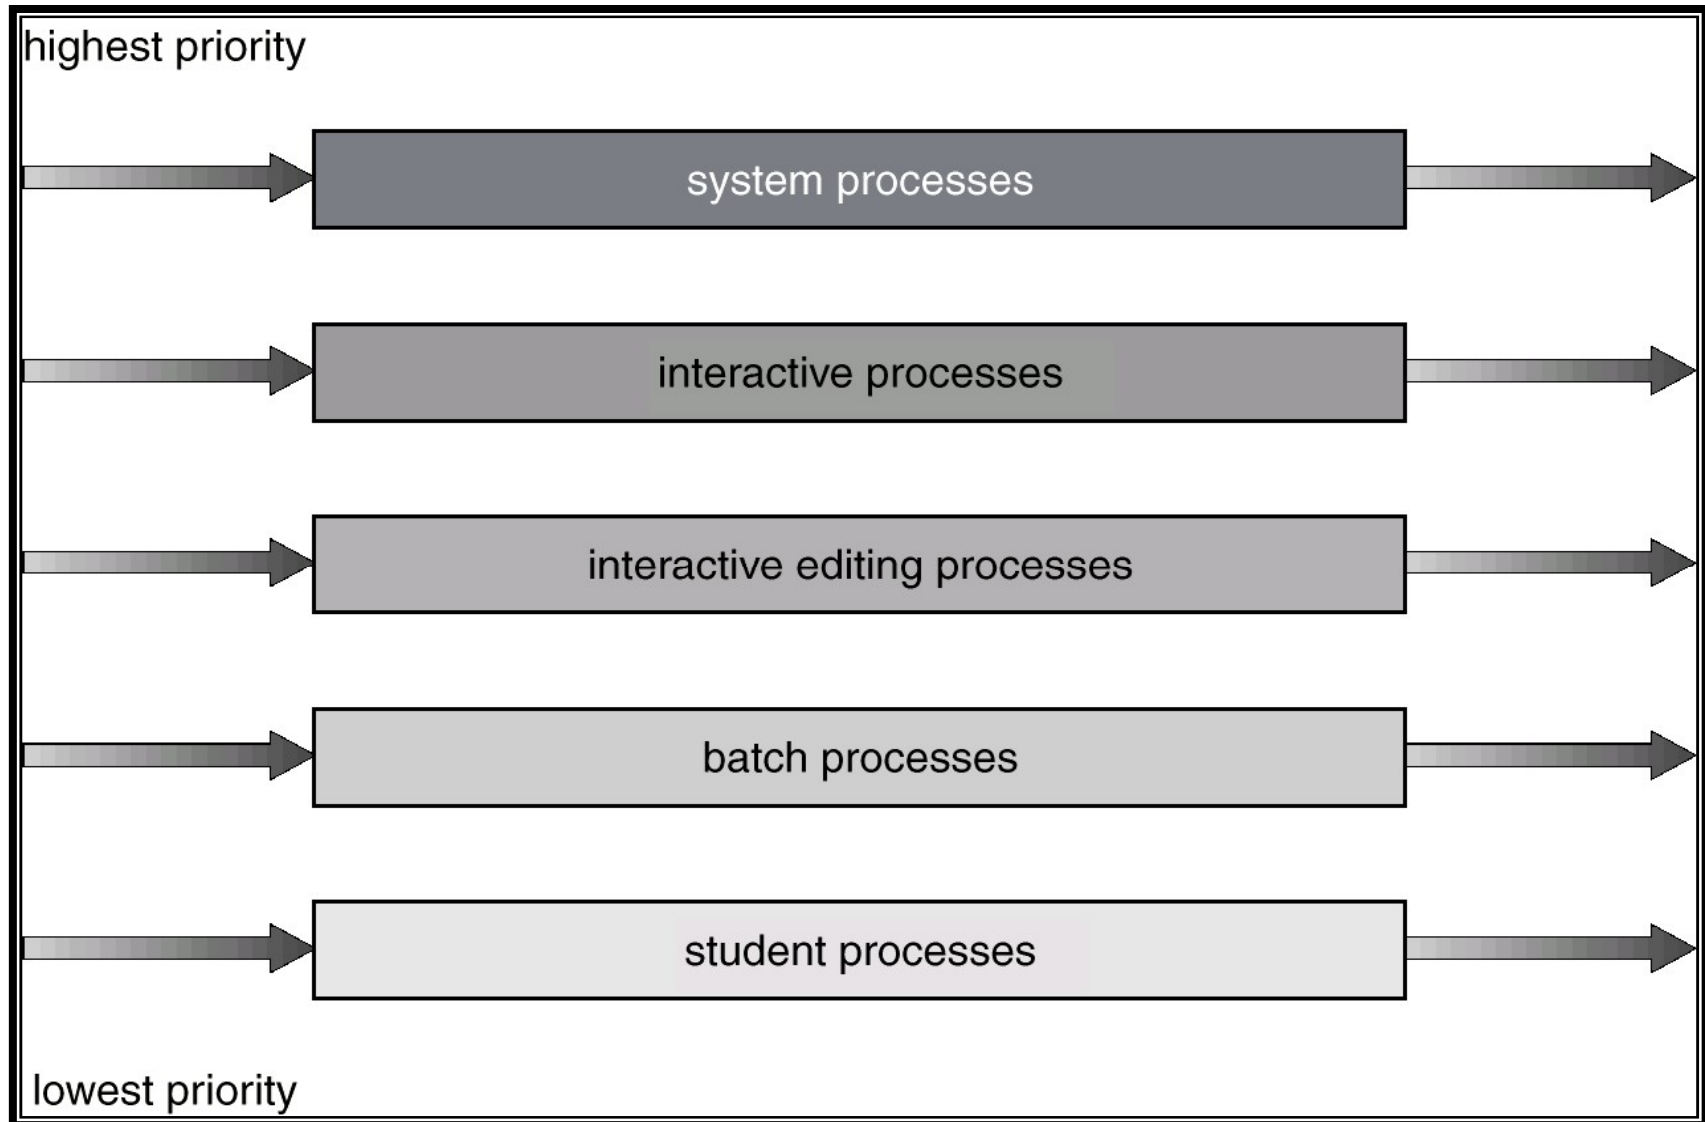
- Increasing the time quantum does not necessarily improve the average turnaround time!



5. Multilevel Queue Scheduling

- Ready queue is partitioned into separate queues:
 - For example, foreground (interactive) and background (batch)
- Each queue has its own scheduling algorithm:
 - Foreground – RR
 - Background – FCFS
- Scheduling must be done between the queues
 - Fixed priority scheduling; (i.e., serve all from foreground then from background)
 - Possibility of starvation
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes
 - i.e., 80% to foreground in RR and 20% to background in FCFS

Multilevel Queue Scheduling



6. Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter
 - when that process needs service


Multilevel Feedback Queue Scheduling

- Example
 - If a process used too much CPU time, then move it to a lower-priority queue
 - If a process waits too long in a lower priority queue, then move it to a higher priority queue

Multilevel Feedback Queue

- Three queues:
 - Q_0 – time quantum 8 milliseconds
 - Q_1 – time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, the job receives 8 milliseconds. If it does not finish in 8 milliseconds, the job is moved to queue Q_1
 - At Q_1 , the job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2

Multilevel Feedback Queue



MLFQ Example

- At $t=0$ there are three CPU bound processes (no I/O) in the system. The CPU burst lengths are: 30 units for P1, 20 units for P2, and 10 units for P3. The system has three RR queues with the following time slices: 1 for queue1, 2 for queue 2, and 4 for queue3. RR- 1 Unit of time slice
- Gantt Chart
- Average Completion Time
- Average Waiting Time

Solution:

The following GANTT chart.

```
p  Process p, p=1,2,3 is running
Qk Queue on level k
```

```

Q1:      123
Q2:      112233          T          T          T
Q3:      111122223333111122223331111222211112222111121111111
Time:      1          2          3          4          5          6
          123456789012345678901234567890123456789012345678901234567890

```

Completion Times

MLFQ P1:60; P2:53; P3:32 Mean Completion: $(60+53+32)/3 = 48 \frac{1}{3}$

Waiting Times

MLFQ P1:30; P2:33; P3:22 Mean Waiting: $(30+33+22)/3 = 28 \frac{1}{3}$

Summary of CPU Scheduling Algorithms

- First-Come, First-Served (FCFS) Scheduling
- Shortest-Job-First (SJF) Scheduling
 - Nonpreemptive
 - Preemptive or Shortest Remaining Time First (SRTF)
- Priority Scheduling
 - Preemptive
 - Nonpreemptive
- Round Robin (RR)
- Multilevel Queue Scheduling
- Multilevel Feedback Queue

References

- Operating System Concepts (Silberschatz, 9th edition)
Chapter 5