

Operating Systems

CS2006


Lecture 14

Memory Management-II

17th April 2023

Dr. Rana Asif Rehman

Paging




Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 1 GB per page
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

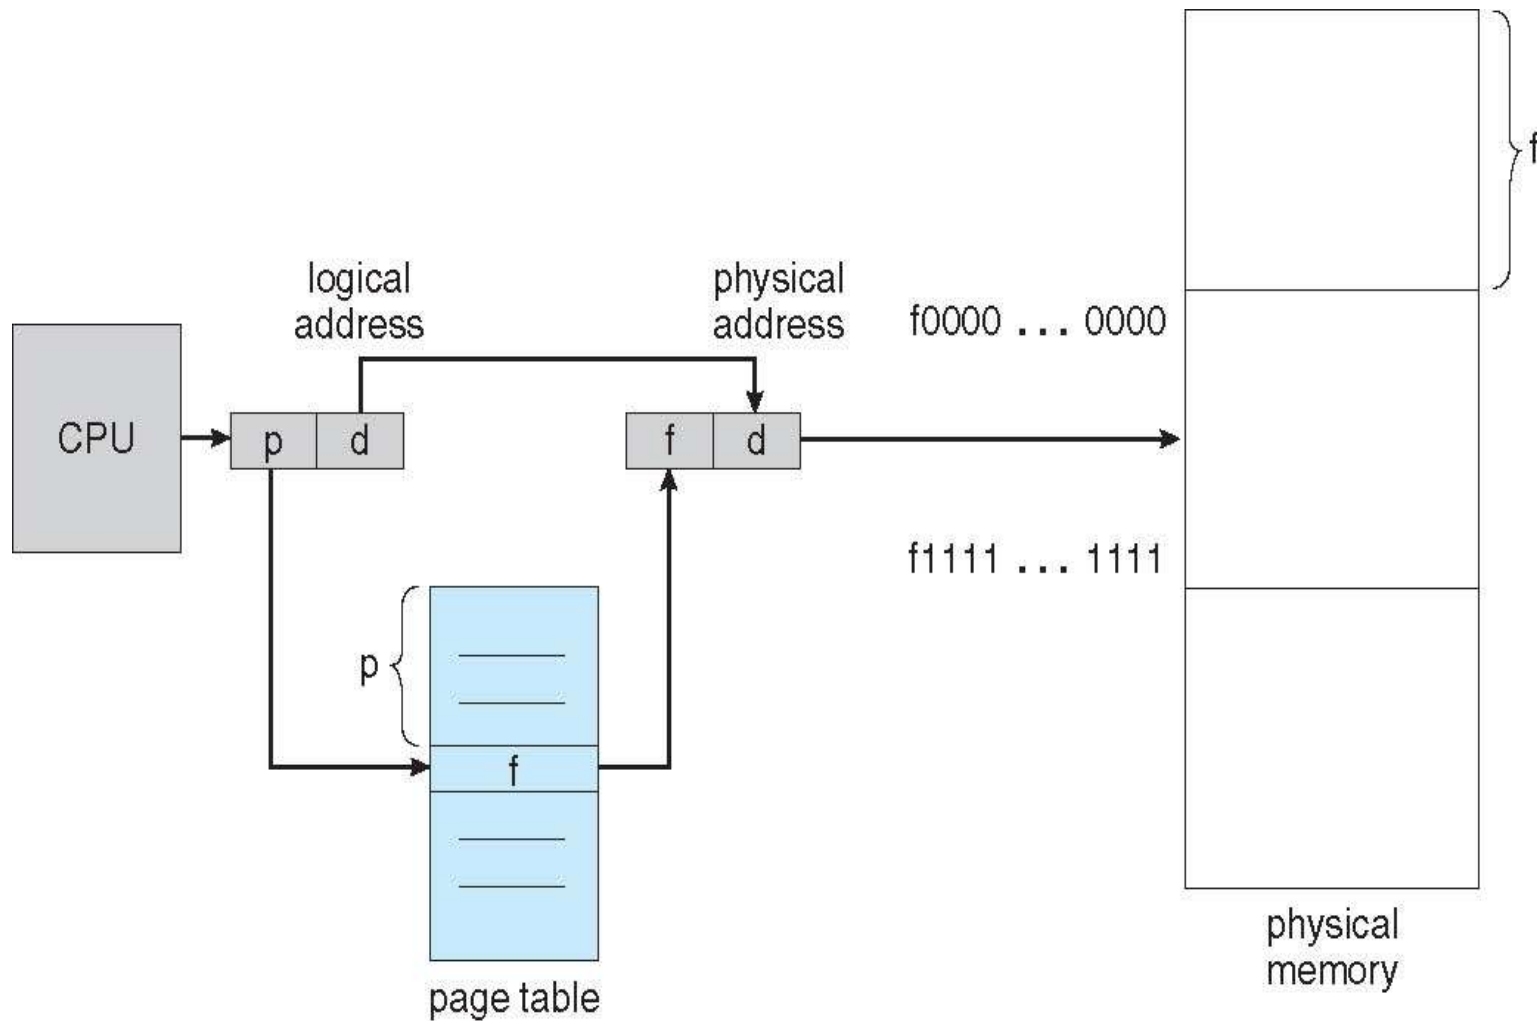
Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit




- For given logical address space 2^m and page size 2^n

Paging Hardware



Paging Model of Logical and Physical Memory



Paging Example

| | |
|----|---|
| 0 | a |
| 1 | b |
| 2 | c |
| 3 | d |
| 4 | e |
| 5 | f |
| 6 | g |
| 7 | h |
| 8 | i |
| 9 | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

| | |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table

| | |
|----|------------------|
| 0 | |
| 4 | i j k l |
| 8 | m n o p |
| 12 | |
| 16 | |
| 20 | a b c d |
| 24 | e f g h |
| 28 | |

physical memory

$n=2$ and $m=4$ 32-byte memory and 4-byte pages

Paging Example

32-byte memory and 4-byte pages

($n=2$ and $m=4$)

- Using a page size of 4 bytes and a physical memory of 32 bytes (8 pages).
- Logical address 0 is page 0, offset 0.
- Page 0 is in frame 5.
- Logical address 0 maps to physical address 20 [= $(5 \times 4) + 0$].
- Logical address 3 (page 0, offset 3) maps to physical address 23 [= $(5 \times 4) + 3$].
- Logical address 4 is page 1, offset 0; according to the page table, page 1 is mapped to frame 6.
- Thus, logical address 4 maps to physical address 24 [= $(6 \times 4) + 0$].
- Logical address 13 maps to physical address 9.

| | |
|----|---|
| 0 | a |
| 1 | b |
| 2 | c |
| 3 | d |
| 4 | e |
| 5 | f |
| 6 | g |
| 7 | h |
| 8 | i |
| 9 | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

| | |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table


| | |
|----|------------------|
| 0 | |
| 4 | i j k l |
| 8 | m n o p |
| 12 | |
| 16 | |
| 20 | a b c d |
| 24 | e f g h |
| 28 | |

physical memory

Paging (Cont.)


- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Free Frames



(a)

Before allocation



(b)

After allocation

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access


Associative Memory

- Associative memory – parallel search

| Page # | Frame # |
|--------|---------|
| | |
| | |
| | |
| | |

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table

00000

| |
|--------|
| page 0 |
| page 1 |
| page 2 |
| page 3 |
| page 4 |
| page 5 |

10,468
12,287

frame number

valid-invalid bit

| | | |
|---|---|---|
| 0 | 2 | v |
| 1 | 3 | v |
| 2 | 4 | v |
| 3 | 7 | v |
| 4 | 8 | v |
| 5 | 9 | v |
| 6 | 0 | i |
| 7 | 0 | i |

page table

| | |
|---|---------------|
| 0 | |
| 1 | |
| 2 | page 0 |
| 3 | page 1 |
| 4 | page 2 |
| 5 | |
| 6 | |
| 7 | page 3 |
| 8 | page 4 |
| 9 | page 5 |
| | ⋮ |
| | page <i>n</i> |

Shared Pages


- **Shared code**

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

- **Private code and data**

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example




Structure of the Page Table

- Memory structures for paging can get huge using straightforward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - If each entry is 4 bytes \rightarrow 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory
- Hierarchical Page Tables
- Hashed Page Tables
- Inverted Page Tables

1. Hierarchical Page Tables


- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table

Two-Level Page-Table Scheme



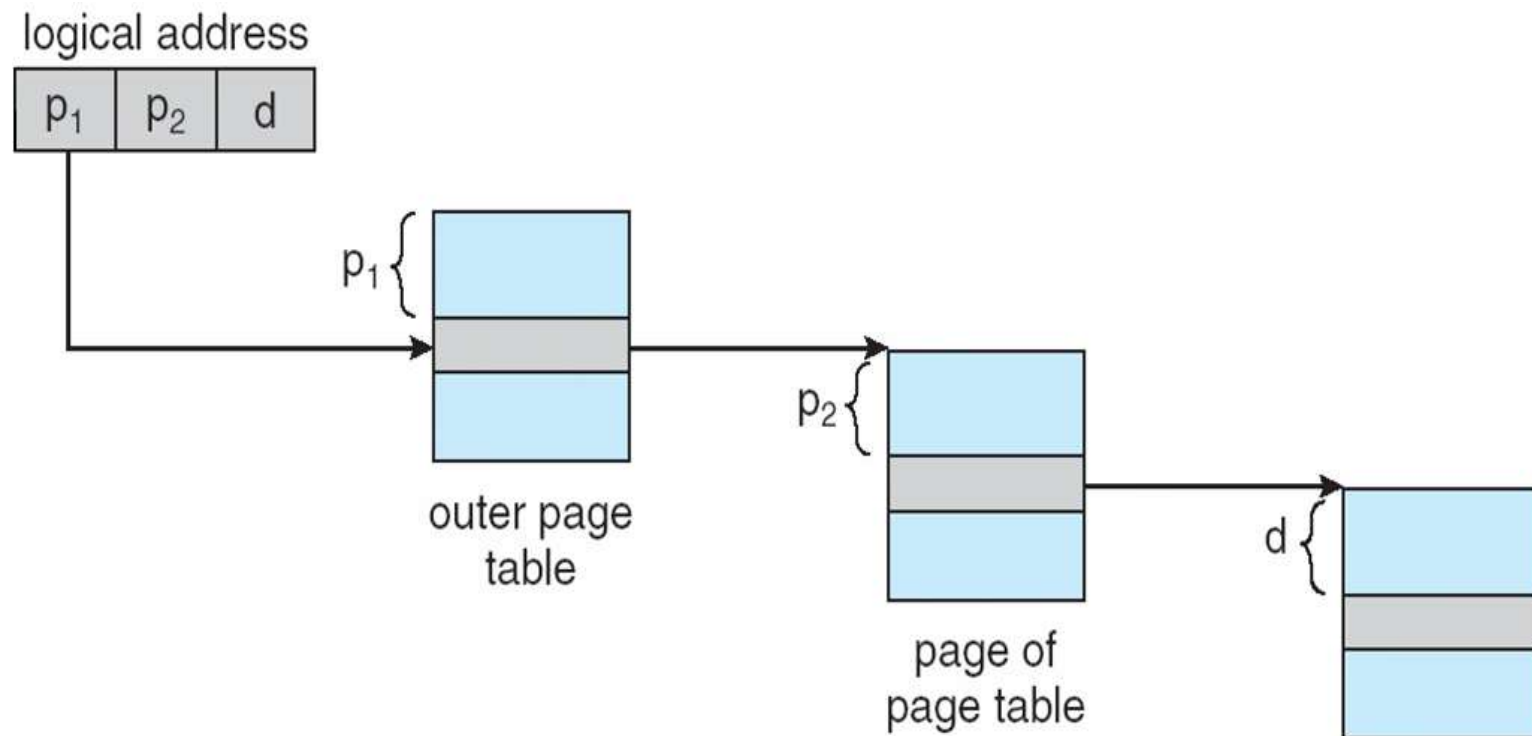
Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:




- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

Address-Translation Scheme





64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB (2^{12})
 - Then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - Address would look like



- Outer page table has 2^{42} entries or 2^{44} bytes
- One solution is to add a 2nd outer page table
- But in the following example the 2nd outer page table is still 2^{34} bytes in size
 - And possibly 4 memory access to get to one physical memory location


Three-level Paging Scheme



2. Hashed Page Table

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)


Hashed Page Table



3. Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
 - TLB can accelerate access
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address

Inverted Page Table Architecture



Numericals

Memory Management

Page Table Size

- The size of a page table is given by the following equation:

- $$\text{Page table size} = \frac{\text{logical address space size} \times \text{page table entry size}}{\text{page size}}$$

- For example:

- Many processors have 32-bit logical addresses, which results in a 4GB logical address space size. The page table entries size is usually 4B. If the page size is 4KB then the page table size is:

- $$\text{Page table size} = \frac{4 \text{ GB} \times 4 \text{ B}}{4 \text{ KB}} = 4 \text{ MB}$$

Paging – Effective Access Time

Find the effective access time when there is an 80% chance of a TLB hit and it takes 100 nanosecond to access the memory.

EAT = Effective Access Time =

$(\text{tlb_hit_ratio} \times \text{memory access time}) +$
 $(\text{tlb_miss_ratio} \times 2 \times \text{memory_access_time})$

$= (0.80 \times 100) + (0.20 \times 200) = 120 \text{ nanoseconds}$

Memory Management – Memory size

- Suppose you have a ram of size 32MB, calculate the **total number of locations/entries** in ram and the **number of bits required to address each location**.

Solution: Supposing the system is byte addressable.

Memory size = 32MB = loc x 1B (where loc is the number of locations available in memory supposing each location is of 1 byte)


$$\text{loc} = 32\text{MB}/1\text{B} = 32\text{M locations} = 32 \times 2^{20} = 2^5 \times 2^{20}$$

(we divide memory size by the size of each location to get # of locations)
 $= 2^{25}$

So the total number of bits required to represent 2^{25} locations is = 25

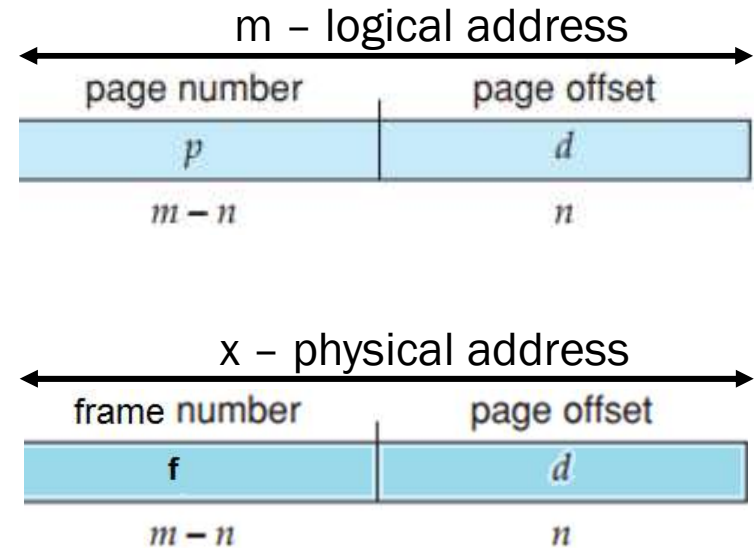
Note: You maybe provided with the size of each location i.e. 1B assumed here.

Systems can be byte or word addressable in general.



Memory Management - Paging

- Supposing logical address = 26 bits
- Physical address = 16 bits
- Page size = 1KB
- **Find the following: # of frames & pages**



Supposing system is Byte addressable. Now consider the following calculations:

Since the page size is 1KB. So, $1\text{KB} = 2^{10} \text{ B}$

So the offset $n = 10$ bits.

Since x is the physical address space, and n is number of offset bits. $x - n =$ physical address location bits. $16 \text{ bits} - 10 \text{ bits} = 6 \text{ bits}$. Hence, 2^6 is the total number of frames.

m is the logical address space. As $n = 10$, $m - n = 26 - 10 = 16$. Hence $2^{16} = 64\text{K}$ pages in the virtual address space

Memory Management - paging

- Consider a logical address space of 64 pages of 1024 words each, mapped onto a physical memory of 32 frames.
 - a. How many bits are required to represent the logical address space?
 - b. How many bits are there in the physical address?

Logical address space bits = m = ?

Recall that the size of the page/frame is = 2^n and the size of logical address space are 2^m

Then, **logical address space = # of pages x page size**

$$\rightarrow 2^m = \# \text{ of pages} \times 2^n$$

$$\rightarrow \# \text{ of pages} = 64 = 2^{m-n}$$

$$n = ?$$

$$\text{Page size} = 2^n \rightarrow 1024 = 2^n$$

$$2^{10} = 2^n \rightarrow n = 10\text{-bits}$$

$$\text{Again: } \# \text{ of pages} = 2^{m-n} \rightarrow 64 = 2^{m-10} \rightarrow 2^6 = 2^{m-10} \rightarrow 6 = m-10 \rightarrow m = 16\text{-bits (logical address)}$$

Physical address: Let x be the physical address bits.

$$\text{Size of physical address space} = 2^x$$

$$\text{Size of physical address space} = \# \text{ of frames} \times \text{frame size (frame size} = \text{page size)}$$

$$\text{Size of physical address space} = 32 \times 1024 \rightarrow 2^x = 2^5 \times 2^{10} \rightarrow 2^x = 2^{15} \rightarrow \text{number of bits} = x = 15 \text{ bits}$$

Memory Management - paging

- Consider a logical address space of 64 pages of 1024 words each, mapped onto a physical memory of 32 frames.
 - a. How many bits are required to represent the logical address space?
 - b. How many bits are there in the physical address?

Easy solution

Addressing within a 1024-words page requires 10 bits because $1024 = 2^{10}$. Since the logical address space consists of 64 pages = 2^6 pages, the logical addresses must be $10+6 = 16$ bits. Similarly, since there are 32 frame = 2^5 physical frames, physical addresses are $5 + 10 = 15$ bits long.

Memory Management - paging

Consider a system using Paging for memory management. The system uses a page size of 1024 bytes, and the length of the address register is 18 bits. Now compute the following parameters:

- a) Size of RAM
- b) Total number of frames in the RAM
- c) Number of bits in a page table entry

Solution:

$$\text{a) } 2^{18} = 2^8 * 2^{10} = 256 * 1024 = 256 \text{ KB}$$

$$\text{b) } 256 \text{ KB} / 1024 \text{ B} = 256$$

$$\text{c) } 18 - 10 = 8 \text{ bits}$$

Memory Management - Paging

Question

Write a C/C++ function to perform address translation in a system using Paging for memory management. You are given with the page size, page table, and a logical block no, and you are required to compute the corresponding physical address. Use the following named constant, and the function prototype:

```
const int PS = ... ;           // page size

int translate(int log,          // logical block no
              int tbl[],        // page table
              int n);           // length of the page table
```

Solution:

```
int translate(int log, int tbl[], int n) {

    int page = log / PS;         // compute page no
    int frame = tbl[page];       // get frame no
    int off = log % PS;          // calculate offset within the page

    return frame * PS + off;     // compute the physical address
}
```

Memory Management - Paging

Consider a system with a page size of 256. Assume a process running in this system has the following page table:

| Page number | Frame number |
|-------------|--------------|
| 0 | 50 |
| 1 | 10 |
| 2 | 90 |

Now translate the following logical addresses into the corresponding paging physical addresses:

i) 700

$$\text{page} = 700 / 256 = 2$$

$$\text{frame} = \text{tbl}[2] = 90$$

$$\text{offset} = 700 \bmod 256 = 188$$

$$\text{phy add} = 90 * 256 + 188$$

$$= 23,228$$

CS-2006 Operating Systems

Memory translation – logical to physical

Given a logical address of 2486 and a page size of 256. What is the corresponding physical address?

We can divide the address by page size to get page #. So $2486/256 = 9.7$
Therefore, the page # is 9. In the page table, we can grab the corresponding frame number.

In order to find the offset, we take the mod of 2486 with 256. So, $2486\%256 = 182$

How to find the exact address?

Supposing that the page table returned 5 as the frame number against the page number 9. The exact memory address of the frame will be:
(locations/entries_Per_Page*frame_Number)-1 (-1 because memory indexing starts at 0)

So $\rightarrow (256 * 5) - 1 = 1279$

Now, we will add the offset to it to get the actual physical memory address:

$1279 + 182 = 1461 = 0x5B5$ is the memory location.

Memory translation – logical to physical(2nd method)

Given a logical address of 2486 and a page size of 256. What is the corresponding physical address?

Convert the logical address to binary number. 2486 in binary is: 100110110110

Now, take the \log_2 of the page size. i.e. $\log_2(256) = \log_2(2^8) \rightarrow 8$

In the above binary number, the least-significant 8 numbers are the page offset. i.e. 1001 10110110
The underlined 0's and 1's is the offset. 10110110 = 182

The most significant left most 4 bits represents the page number.

offset
↓
1001 10110110
↑
page number

Page number is therefore 9.

How to find the exact address?

Supposing that the page table returned 5 as the frame number against the page number 9. The exact memory address of the frame will be: $(\text{locations/entries_Per_Page} * \text{frame_Number}) - 1$ (-1 because memory indexing starts at 0)

So $\rightarrow (256 * 5) - 1 = 1279$

Now, we will add the offset to it to get the actual physical memory address: $1279 + 182 = 1461 = 0x5B5$
Is the memory location.

Paging guidance

Paging arithmetic laws:

Page size = frame size

Logical address space (/size) = 2^m

Logical address space (/size) = # of pages \times page size

Physical address space (/size) = 2^x (*where x is the number of bits in physical address*)

Physical address space (/size) = # of frames \times frame size

Page size = frame size = 2^n

of pages = 2^{m-n}

of entries (records) in page table = # of pages

References

- Chapter 3, Modern Operating System
- http://cseweb.ucsd.edu/classes/sp00/cse120_A/mem.html
- Operating System Concepts (Silberschatz, 9th edition)
Chapter 8